

MESSAGEFOUNDRY DOCUMENTATION

Adoption guide

This guide is for teams piloting MessageFoundry (MEFOR) — an open-source, Python healthcare integration engine — and taking it from a first install to full production use.

This guide is for teams piloting **MessageFoundry (MEFOR)** — an open-source, Python healthcare integration engine — and taking it from a first install to full production use. It is an **orchestration** document: it ties the existing docs together and adds the install-to-production **rollout plan** that nothing else covers. Where a topic has a dedicated reference, this guide links to it rather than repeating it.

Read this section first. MessageFoundry is **pre-1.0** software. It does a small set of things well and production-grade today (see §2), and it has clearly-bounded areas that are **experimental or not yet built**. Any new integration engine — this one included — *will* have bugs you have not hit yet. The whole point of the staged rollout in §11 is to find them where they are cheap (a lab, a shadow feed) instead of where they are expensive (a production cut-over). If you adopt MEFOR, adopt the rollout discipline with it.

Table of contents

1. [What MessageFoundry is, and who should pilot it](#)
2. [Maturity & honest limitations — read before you plan](#)
3. [Prerequisites & environment checklist](#)
4. [Installation](#)
5. [Minimum viable configuration](#)
6. [Security & PHI hardening before real data](#)
7. [Reliability configuration — how nothing gets lost](#)
8. [Pre-traffic validation](#)
9. [Capacity & load testing on *your* hardware](#)
10. [Backup, restore & disaster recovery](#)
11. [Staged rollout plan with go/no-go gates](#)
12. [Day-2 operations & monitoring](#)
13. [Upgrade & rollback](#)
14. [High availability — rollout & VIP standup runbook](#)
15. [Getting help & reporting bugs](#)
16. [Decommissioning a pilot](#)

1. What MessageFoundry is, and who should pilot it

MessageFoundry routes, transforms, and validates messages between **Connections** — HL7 v2.x by default, with payload-agnostic support for other formats — and its routing and handling are written in **Python**. The runtime model is a graph wired by name:

- **Connection** — an endpoint that receives (inbound) or sends (outbound) messages (MLLP, TCP, File today; REST/SOAP/Database destinations and a Database poll source also ship — see [CONNECTIONS.md](#)).

- **Router** — a Python function bound to an inbound connection that decides which Handler(s) see each message.
- **Handler** — a Python function that filters → transforms a message and emits `Send`s to outbound connections.

The engine is a **headless asyncio service** (FastAPI/uvicorn) that owns a durable message store and supervises one worker set per connection. A browser **web console** (`/ui`) and a **VS Code extension** operate it over a localhost HTTP API. See [ARCHITECTURE.md](#) for the full model.

Who should pilot it now. Teams who want a Python-native, open-source alternative to Mirth/Corepoint and who are comfortable validating a pre-1.0 tool against their own traffic before trusting it. A single engine node on a trusted network is the simplest pilot; **native TLS** (API + MLLP) and an opt-in **active-passive failover** cluster on a shared server-DB store (PostgreSQL or SQL Server) are both built when you need them (see §2/§6/§14). What is genuinely *not* there yet is MFA, off-box log shipping, and a de-identification framework — track those items (§2) and pilot the parts that are ready. (Horizontal *active-active* scale-out was dropped and is not a planned milestone; active-passive HA is the supported HA model.)

2. Maturity & honest limitations — read before you plan

MessageFoundry is solid for single-node production, and now also supports **opt-in active-passive failover** (a leader/standby cluster on a shared server-DB store — PostgreSQL or SQL Server; §14). The authoritative built-vs-roadmap references are [ARCHITECTURE.md](#) and the README "Roadmap" section — use the table below alongside them when planning.

Built and production-ready

Capability	Status
Code-first Connection/Router/Handler graph	✓ Built
SQLite (WAL) store backend	✓ Production-ready — the default, single-node/dev
PostgreSQL store backend (single-node)	✓ Production-ready — full staged pipeline, at-rest encryption, retention; single-node parity with SQLite
Microsoft SQL Server store backend (single-node)	✓ Production-ready — full staged pipeline + response capture, at-rest encryption; needs the <code>sqlserver</code> extra + OS-level ODBC Driver 18. Engine-enforced retention/purge at full parity with SQLite; only WAL checkpoint / <code>VACUUM</code> / the DB-tier <code>.mfbak</code> snapshot are DBA-owned there.
Transactional staged queue (ingress→routed→outbound), at-least-once, dead-letter, replay	✓ Built — see ADR 0001
Auth + RBAC + hash-chained audit log	✓ Built — see SECURITY.md
At-rest body encryption (AES-256-GCM, opt-in) + key rotation	✓ Built — see PHI.md
MLLP / TCP / File connectors; REST / SOAP / Database destinations; Database poll source	✓ Built — see CONNECTIONS.md
Validation & load tooling (<code>generate</code> , <code>check</code> , <code>dryrun</code> , the test harness, the load harness)	✓ Built — see §8/§9 and LOAD-TESTING.md
Windows-service deployment via NSSM	✓ Built — see SERVICE.md
Native transport TLS (API + MLLP)	✓ Built — in-process API TLS (HTTPS/WSS) + per-connection MLLP-over-TLS, ≥TLS 1.2, opt-in mTLS, and a fail-closed off-loopback bind guard (a non-loopback bind without TLS is refused). Raw TCP/X12 stay plaintext (loopback/proxy). See DEPLOYMENT.md .
Native MFA (TOTP, local accounts)	✓ Built — RFC 6238 TOTP + single-use recovery codes; <code>[auth].require_mfa</code> enforces a second factor for local Administrators at the step-up boundary. AD/Entra users' MFA stays delegated to the IdP. See SECURITY.md .
Off-box log + audit forwarding	✓ Built — <code>[logging].forward_*</code> ships operational logs + PHI-redacted audit rows to a syslog/SIEM collector, over native TLS when you set <code>forward_protocol = "tls"</code> (RFC 5425, ADR 0080; port 6514, CA anchor via <code>forward_tls_*</code>). Residual: the transport default is UDP, so TLS is a per-

Capability	Status
	deployment opt-in — set it, or front the collector with a local TLS-forwarding agent. See PHI.md §7.
Active-passive HA / failover	✅ Built (Track B) — opt-in leader/standby cluster on a shared server-DB store (PostgreSQL or SQL Server): only the leader runs the graph, self-fencing leadership lease, immediate on-promotion recovery. Single-node stays the byte-identical default. See CLUSTERING.md + §14.

Experimental or not yet built — do not depend on these for a production pilot

Capability	Status & implication
Transport TLS for raw TCP / X12	❌ Not built — those two connectors are plaintext-only; keep them on loopback or front with a TLS-terminating proxy. (API + MLLP do have native TLS — see §6/ DEPLOYMENT.md .)
<code>ack_after=delivered</code> (defer the ACK until downstream delivery)	❌ Not built — requesting it is rejected at config load. Only ACK-on-receipt exists, so a routing/transform/delivery failure happens after the sender was already told <code>AA</code> and will not NAK back. Operators rely on the message disposition + alerts, not the ACK.
De-identification framework	❌ Not built. The AI assistant's <code>deidentified</code> scope falls back to <code>code_only</code> .
In-place SQLite → server-DB migration	❌ Not built. Server-DB deployments are greenfield only — there is no automatic carry-over of SQLite history. Drain and cut over deliberately (§13).
A throughput guarantee for your hardware	⚠️ By design. A baseline + tuning method is published (TUNING-BASELINE.md , Gate #3) as a two-tier gate — host-independent conformance invariants (hard) + performance numbers " <i>as measured on the reference config</i> ". Because the durable-write path is hardware-dependent, those msg/s are not a promise for your box. Measure on your own hardware (§9).

The early-adopter bargain, stated plainly: you get a durable engine with native TLS, real auth, opt-in active-passive failover, and a real validation toolchain, in exchange for validating capacity on your own hardware and supplying the operational pieces that aren't built yet (MFA, off-box logs, de-identification). If that trade is acceptable, the rest of this guide is your playbook.

3. Prerequisites & environment checklist

Consolidate these before you install anything:

- [] **Python 3.14+** on the engine host.
- [] **OS:** Windows is the primary supported/serviced platform (NSSM); the engine itself is cross-platform Python.
- [] **Administrator/elevation** on the host if you will install the Windows service.

- [] **Outbound network access** for the service installer to download the SHA-256-pinned NSSM binary (or pre-stage NSSM on the host / on `PATH`).
- [] **Firewall plan**: open your inbound MLLP listener port(s) (the samples use e.g. `2575` / `2600`) to senders, and decide who may reach the **API on** `127.0.0.1:8765` (default loopback — keep it that way; see §6).
- [] **A writable data directory** for the store + logs (service default: `C:\ProgramData\MessageFoundry`).
- [] **Backend decision (made here, not later): SQLite** (default, zero extra deps) for a single-node pilot, or a server DB — **PostgreSQL** (`messagefoundry[postgres]`, pure-Python) or **SQL Server** (`messagefoundry[sqlserver]` + OS-level ODBC Driver 18) — if you want a server store or a path toward DB-tier HA. See §2.
- [] If you will run the **cluster** path (active-passive failover): **NTP time sync** across nodes is a hard prerequisite, every node needs the **same config dir**, and `[store].backend = "postgres"` or `"sqlserver"`. (Single-node pilots skip this entirely; see §14.)
- [] A **PHI encryption key** plan (§6) and a **backup target + key-escrow** plan (§10) decided before any real data flows.

4. Installation

New here? Start with the [Installation Guide](#) — the focused walkthrough of installing the engine and standing up your own private **config repo**, including running multiple instances from one repo. This section is the rollout-oriented summary of the same material.

Full reference: the [Installation Guide](#) and [SERVICE.md](#). The essentials:

4.1 Install the engine

MessageFoundry is a **read-only, version-pinned dependency** ([ADR 0017](#)): install a published wheel and **pin the exact version**, the same way you pin any other production dependency. Create a venv and install:

```
python -m venv .venv
.\.venv\Scripts\Activate.ps1
pip install "messagefoundry==0.1.0"           # pin the exact engine version (core runtime only)
```

`messagefoundry==0.1.0` pulls only the **core runtime** — what a headless engine needs. Add extras (§4.2) for the PySide6 test harness, a server-DB backend, or SFTP; the browser web console installs as its own `messagefoundry-webconsole` wheel.

⚠ Early access. `0.1.0` is an **Early Access** release on public PyPI — feature-complete and test-validated, but the external review + pen test that gate a security-certified **v1.0** land after launch. The exact-pin command above (`==0.1.0`) resolves today; the earlier `0.1.0rc1` pre-release also remains installable. You can equally install from the engine's **GitHub Release assets** or your organization's **private index**.

Verify the release before you install. MessageFoundry ships one signed wheel to many PHI-bearing instances, so verify the artifact's provenance *before* installing — pinning a version (or a hash) proves you got a *fixed* file, not that it is the one MessageFoundry built. Every release carries **SLSA build provenance** and a **Sigstore signature**; check both with the **GitHub CLI** (`gh` ≥ 2.49), and optionally `sigstore` (`pip install sigstore`). Install **only** the file that passes:

```
$V = "0.1.0" # the exact version you intend to install

# Download the wheel + its Sigstore bundle from that release's assets
gh release download "v$V" --repo MEFORORG/MessageFoundry `
  --pattern "messagefoundry-$V-*.whl" --pattern "messagefoundry-$V-*.whl.sigstore*"

# Verify SLSA build provenance: artifact -> source commit -> builder workflow
gh attestation verify "messagefoundry-$V-py3-none-any.whl" --repo MEFORORG/MessageFoundry

# (defense in depth) Verify the Sigstore signature pins the release workflow identity
python -m sigstore verify identity "messagefoundry-$V-py3-none-any.whl" `
  --cert-identity
  "https://github.com/MEFORORG/MessageFoundry/.github/workflows/release.yml@refs/tags/v$V" `
  --cert-oidc-issuer "https://token.actions.githubusercontent.com"

# Only if BOTH pass, install the exact file you verified
pip install ".\messagefoundry-$V-py3-none-any.whl"
```

The same attestation also covers the **public PyPI** copy (it is byte-identical), so you can `pip download "messagefoundry==$V" --no-deps -d .\verify`, `gh attestation verify` the downloaded wheel, then `pip install --no-index --find-links .\verify "messagefoundry==$V"`. A registry/mirror substitution or a relabelled file **fails** the check. (The `--cert-identity` ref must match the tag you install — e.g. `refs/tags/v0.1.0-rc1` for a pre-release.)

For a **reproducible pinned** deploy, generate a hash-locked requirements file scoped to the extras you actually run and install it with `--require-hashes`. The scaffolded config repo (`messagefoundry init`, below) already pins the engine in its `requirements.txt` — extend that into a full hash-lock for your host.

4.2 Optional extras

Extra	Pulls in	When
postgres	asyncpg (no OS dep; ships compiled wheels)	Using the PostgreSQL backend (recommended prod path)
console	PySide6 + keyring	Running the desktop admin console
sftp	paramiko	SFTP connectors
sqlserver	aiodbc + OS-level Microsoft ODBC Driver 18	The SQL Server store backend (backend=sqlserver , production) and the DATABASE connector family.
dev	pytest/ruff/mypy/httpx	Development & CI

⚠ There is **no friendly preflight** for the `postgres` extra: if you set `backend=postgres` but forgot `pip install 'messagefoundry[postgres]'`, you get a raw `ImportError` at startup instead of a clear message. Install the extra with the backend.

Start your own config repo (`messagefoundry init`)

§4.1 installed the **engine**. Now scaffold the other half a deploying organization owns — your **own** separately-versioned **config repo** ([ADR 0017](#)) — which holds your Connections/Routers/Handlers and drives one or more engine instances. **This is the recommended way to run MessageFoundry:** the `samples/` directory used in older quickstarts ships only in a source checkout, **not in the installed wheel**, so a wheel install runs against *your* `config/`, not `samples/`.

```
messagefoundry init ./my-config-repo
```

It writes a runnable starter feed (`config/`), `environments/<env>.toml` value stubs, a synthetic fixture, an instance `messagefoundry.toml` (active environment + posture), a `requirements.txt` pinning this engine version, a CI `check` workflow, and `.vscode` settings — so `messagefoundry check --config config --messages messages/sets` is green from the first commit. See the generated `README.md` for the day-to-day workflow.

4.3 Run it (foreground, to learn the ropes)

From your config repo (the one `messagefoundry init` created), run the engine against its `config/` directory:

```
cd ./my-config-repo
python -m messagefoundry serve --config config --db ./messagefoundry.db --env dev
```

`serve` flags and their precedence (**CLI** > `MEFOR_<SECTION>_<KEY>` `env` > `messagefoundry.toml` > **built-in default**): `--config` (your graph directory — pass `--config config` for a scaffolded repo; the built-in default `samples/config` exists only in a source checkout), `--service-config` (default `./messagefoundry.toml` if present), `--db`, `--host`, `--port`, `--log-level`, `--env` (a **free-form** environment name, [ADR 0017](#)), `--allow-insecure-bind`.

⚠️ **The active environment is required.** `serve` refuses to start (exit 2) without `--env <name>` (or `[ai].environment`) — there is no silent `prod` default, so a missing env can never resolve another environment's values/secrets. Built-in names `dev` / `staging` / `prod` carry a default posture; a custom name (e.g. `test`, `poc`) also needs `[ai].data_class` + `[ai].production`. The active environment is logged at startup.

🔑 `dev` now carries the PHI posture (**ADR 0148**) — provide a store key or declare synthetic. Since ADR 0148 (GIVEN 1) the built-in `dev` env derives the **PHI** data-class, so your first run exercises the same at-rest-encryption path production uses (rather than first meeting it in prod). `serve --env dev` therefore **refuses to start (exit 2) without a store encryption key**. Two ways forward for a local run: - **Recommended — mint a throwaway dev key** (exercises the real encryption path): run `messagefoundry gen-key` and set the printed base64 value as `MEFOR_STORE_ENCRYPTION_KEY` (a dev key is fine; **never commit it**). - **Genuinely no-PHI box** — declare it synthetic: set `[security].handles_real_patient_data = false` (a loud, audited opt-out) to run **key-free**, for a dev/CI box that only ever processes synthetic HL7.

The refuse/warn severity of the PHI serve-gate ladder is the `[security].enforcement` dial (default `enforce`, byte-identical to the former production behaviour). On a **loopback** dev bind you hit only the keyless-PHI refusal above — the off-loopback exposure rungs (TLS, MFA-at-exposure, ...) need a non-loopback bind. Set `[security].enforcement = warn` to downgrade the PHI refusals to loud, audited warnings during local bring-up.

4.4 Run it as a Windows service (the supported production run-mode)

Use the elevated installer; install under a **least-privilege virtual account** rather than the default LocalSystem:

```
# from an elevated shell
scripts\service\install-service.ps1 -ServiceAccount "NT SERVICE\MessageFoundry"
```

The installer is idempotent, auto-downloads a SHA-256-pinned NSSM, bakes absolute `serve` paths into the service, and (with `-ServiceAccount`) auto-grants config-read + data-dir-read/write to the account. Service defaults: name `MessageFoundry`, data dir `C:\ProgramData\MessageFoundry`, store `<DataDir>\messagefoundry.db`, logs `<DataDir>\logs`, bind `127.0.0.1:8765`.

⚠️ **Pinned-wheel operational model.** With a pinned-version install (§4.1), the running service loads the **installed wheel** — a known, pinned version, not a moving checkout. Picking up a new engine version is a deliberate `pip install "messagefoundry==<new>"` + NSSM restart (§13), so every upgrade is an explicit, reviewable act. (A contributor running the editable install instead serves whatever branch is checked out — treat that checkout as the release artifact; see §13.)

4.5 First-run admin bootstrap

Auth is **enabled by default**. On the first start against an empty store, MEFOR creates a one-time bootstrap admin (`admin`) and writes its password to an **owner-only** `bootstrap-admin.txt` next to the store (only the

file *location* is logged — never the password). Then:

1. Log in as `admin` ; you are **forced to change the password** on first use.
2. **Create a second real administrator** promptly.
3. **Delete** `bootstrap-admin.txt` .

The bootstrap account auto-retires once a second admin exists, or — while still unclaimed — 72h after creation.

4.6 Verify it runs

```
curl http://127.0.0.1:8765/health          # -> {"status":"ok"}
# tail <DataDir>\logs\service.out.log for the "wiring started" banner
```

Then send a synthetic message to confirm the end-to-end path. The scaffolded starter feed listens on MLLP `2575` and ships a PHI-free fixture at `messages/sets/example_adt.hl7` — send it with any MLLP client. The convenience senders (`samples/send_mllp.py` , `python -m harness`) ship with the **engine source checkout**, not the installed wheel; from a checkout you can run:

```
python samples/send_mllp.py samples/messages/adt_a01.hl7
```

If start fails, check `service.err.log` first — the common causes are relative paths resolving to the system dir under a service account, a busy MLLP/API port, or a data dir the account can't write.

5. Minimum viable configuration

Full reference: [CONFIGURATION.md](#) (service settings) and [CONNECTIONS.md](#) (the graph).

There are two distinct configuration surfaces:

1. **The message graph (Python modules)** in your `--config` directory. The minimum first flow is one module: an `inbound()` with a transport spec and a `router=` binding, a `@router` that returns handler name(s), and a `@handler` that returns `Send(...)` to a declared `outbound()` . The scaffolded repo (`messagefoundry init` , §4) gives you a working `config/IB_EXAMPLE_ADT.py` to start from (or, from a source checkout, copy `samples/config/IB_ACME_ADT.py`). The loader globs `*.py` (non-recursive; skips `_*` -prefixed helper files), then merges an optional `connections.toml` .
2. **Service/operational settings** in `messagefoundry.toml` (+ `MEFOR_*` env + CLI). Keep **all secrets out of this file and out of source control** — supply them via `MEFOR_<SECTION>_<KEY>` env vars (the loader *warns* if it sees a known secret in the file).

Guidance for a clean first flow:

- **Use the MLLP/File pair** for an initial end-to-end test — both are fully built and need no extras. The Database connector family is production-supported but adds the `[sqlserver]` extra + ODBC Driver 18; MLLP/File keep the first hop dependency-free.

- **Never set a host on an inbound MLLP/TCP connection** (it is a config error). Set the listen interface once, service-side, via `[inbound].bind_host` (loopback for dev; a specific NIC behind a firewall for prod). Outbound MLLP/TCP *do* take the downstream host.
- **Author anything environment-specific as `env("key")`**, put non-secret values in `environments/dev.toml` / `environments/prod.toml` with identical keys, and inject secrets only via `MEFOR_VALUE_<KEY>`. A referenced-but-undefined key fails loud at load. Use `current_environment()` (not `env()`) inside a handler to branch on the deployment.
- **`connections.toml` (data) is optional (ADR 0007)**: move *transport config* there if you want hand/GUI editing; keep *logic* (routers/handlers) in `.py`. A name declared in both a module and `connections.toml` is a hard error (no silent shadowing).
- **The `--config` directory is a trust boundary**. `serve` and `POST /config/reload` **execute** the Python in it, in-process, as the service account. On POSIX the loader refuses a group/world-writable config dir; **on Windows this is your responsibility** — lock the directory's ACL to admins + the service account.

6. Security & PHI hardening before real data

Full references: [SECURITY.md](#), [PHI.md](#), and [DEPLOYMENT.md](#) (network exposure). MEFOR ships real auth, RBAC, audit, opt-in at-rest encryption, and **native TLS** (API + MLLP, with a fail-closed off-loopback bind guard); the remaining transport gaps are **MFA** and **off-box log shipping**. Complete this checklist **before any real PHI flows**:

- [] **API off-loopback requires native TLS**. The API binds `127.0.0.1` by default. To reach it from another host, configure **in-process TLS** (`[api].tls_cert_file` + `[api].tls_key_file`, `tls_min_version` $\geq$ 1.2, opt-in mTLS via `tls_client_ca_file`) **or** front it with a TLS terminator (`[api].tls_terminated_upstream = true` + `[api].trusted_proxies`). A non-loopback bind **without** TLS (or a trusted terminator) is **refused at startup**. **Never use `--allow-insecure-bind` for real PHI** — it is a loud dev-only escape that puts bearer tokens and PHI on the wire in cleartext. (With auth disabled, a non-loopback bind is refused unconditionally.)
- [] **MLLP off-loopback requires native TLS too**. MLLP-over-TLS is built: set `tls = true` + `tls_cert_file` / `tls_key_file` per connection (opt-in mTLS via `tls_ca_file`; $\geq$ TLS 1.2). MLLP is **plaintext by default**, and a non-loopback plaintext MLLP bind is refused. **Raw TCP and X12 have no transport TLS** — keep them on a trusted segment or proxy-terminate. Full matrix: [DEPLOYMENT.md](#).
- [] **Turn on at-rest encryption and make it mandatory**: mint a key with `messagefoundry gen-key` (or a Windows DPAPI-protected key file via `messagefoundry protect-key`), set `MEFOR_STORE_ENCRYPTION_KEY`, and set `[store].require_encryption = true` so the engine refuses to start unencrypted.
- [] **Enable volume encryption (BitLocker/LUKS)**. App-level encryption protects message *bodies*; the `summary` / `control_id` / `message_type` columns and the `-wal` / `-shm` /temp files are **not** app-encrypted and rely on volume encryption.

- [] **Run under a least-privilege account** (the virtual account from §4.4) and lock down the store directory and any File-connector spill directories. **Treat backups as PHI.**
- [] **Finish the bootstrap-admin handoff** (§4.5): change the password, create a second admin, delete `bootstrap-admin.txt`.
- [] **For Active Directory:** use **LDAPS** with a trusted CA, never set `MEFOR_ALLOW_INSECURE_TLS` in production, and configure the directory's lockout/complexity policy (the engine's account lockout covers local accounts only). AD/Entra MFA is enforced by your directory; **local accounts** use the engine's **native TOTP MFA** (`[auth].require_mfa` , WP-14) — enable it before an off-loopback PHI exposure.
- [] **Populate the fail-closed** `[egress]` **allowlist** (it defaults to unrestricted) for REST/Database destinations.
- [] **Keep logging at** `INFO` **or above** and `expose_docs` off in production. Full payloads are never logged at INFO+ by design, but PHI-log-redaction of chained-exception traceback text is not yet fully closed — **do not raise the service to DEBUG with real PHI.**
- [] **Author routers/handlers so they never interpolate raw HL7 into an exception message** (it can surface in `last_error` / `detail`).
- [] Run `messagefoundry audit-verify` periodically (the audit log is *tamper-evident*, not *tamper-proof*), and set `[retention]` windows — they are **off by default (kept forever)**.

7. Reliability configuration — how nothing gets lost

This is the heart of operating a new tool safely. The durability model is a **transactional staged queue** (no external broker): each message flows ingress → routed → outbound, with every handoff a single committed transaction, giving **at-least-once** delivery with crash-safe re-runs. Details in [ADR 0001](#).

Key semantics to internalize:

- **ACK-on-receipt.** The sender is `AA` 'd as soon as the raw message is durably committed (after synchronous decode/parse/optional strict-validate, which still NAK). **Any routing/transform/delivery failure happens after the ACK** and surfaces as an internal **disposition + alert**, never a NAK. Operators monitor disposition + alerts, **not** the ACK, for post-ingress failures.
- **Disposition lifecycle:** `RECEIVED` → `ROUTED` / `UNROUTED` → `PROCESSED` / `FILTERED` / `ERROR` . The store finalizer is the **sole authority** and never finalizes while any stage row is still in flight. Note: a single dead row at *any* stage flips the whole message to `ERROR` **even if a sibling handler delivered** — so read the **per-message event trail**, not just the headline status.
- **Failure classification & policy (per outbound):**
 - Permanent partner reject (`AR` / `CR`) → **dead-letter immediately** (still replayable).
 - Transient (`AE` / `CE`) or transport error → **retry per** `RetryPolicy` .
 - Internal/code error → either **STOP** the lane and raise a `connection_stopped` alert, or **CONTINUE** (auto-dead-letter the bad message and keep flowing).

- `RetryPolicy.max_attempts` **unset = retry forever** (nothing silently lost) with exponential backoff. Under the default **FIFO** ordering, a permanently-failing head **blocks its lane** until it succeeds or is purged.

Mandatory before go-live:

- [] **Wire real alerts.** Configure the `[alerts]` **webhook and/or email** notifier — do **not** rely on the default logging-only sink. The conservative defaults (FIFO head-of-line blocking, retry-forever, STOP-on-internal-error) are only safe if a human gets paged when a lane stalls.
- [] **Set `[delivery]` buildup thresholds** (`max_oldest_seconds` defaults to 300s; set a `max_depth` sized to each connection's throughput) so `queue_buildup` fires before a stuck lane silently backs up. Buildup detection now covers the ingress and routed stages too, not just outbound.
- [] **Choose `RetryPolicy` per outbound deliberately:** retry-forever for partners that must never lose a message (accept head-of-line blocking + rely on buildup alerts), or a finite `max_attempts` where stale data is worse than a replayable dead-letter.
- [] **Choose `InternalErrorPolicy` intentionally:** `CONTINUE` (default) for high-volume feeds where uptime matters most; `STOP` for low-volume feeds where ordering/no-loss matters more than uptime.
- [] **Code routers/handlers as pure and idempotent.** At-least-once means a message can re-run after a crash or a replay. No side-effecting writes mid-transform; the **one** allowed exception is a **live, read-only DB lookup**. Downstream connectors/partners must **dedupe** (e.g. on MSH control id).

Recovery tools you should know cold: `/dead-letters` (triage) + `/dead-letters/replay` (bulk *outbound* recovery), and per-message `/messages/{id}/replay` (for dead ingress/routed rows — router/transform errors, undecryptable raw, a removed handler). Startup automatically returns stale in-flight rows to pending (crash recovery) and dead-letters rows whose destination/handler left the config.

8. Pre-traffic validation

Prove correctness **before** any network traffic. None of this should ever run against real PHI — `generate` / `dryrun` can emit full message bodies; never redirect their output to a committed file or CI log.

1. **Build a synthetic corpus:** `messagefoundry generate --type ADT --count 50 --out <fixtures>` (conformant HL7 v2.5.1, validated against hl7apy; 13 message types, 57 ADT triggers; PHI-free).
2. **Gate the config in CI / a pre-commit hook:** `messagefoundry check --config <dir> --messages <fixtures>`. - `validate` (every module loads; every inbound→router reference resolves; no port collisions) is **required and blocking**. - `dryrun` is **required only when you supply a fixtures dir containing `*.hl7`** — **without fixtures the dryrun is silently skipped** and the gate passes on `validate` alone, so a transform that errors at runtime is *not* caught. **Build and maintain the fixtures.** - `ruff` / `mypy` are advisory (never block).
3. **Inspect the wiring:** `messagefoundry validate --json` (all problems at once) and `messagefoundry graph --config <dir>` (confirm the wired graph matches intent).
4. **Confirm dispositions:** `messagefoundry dryrun` runs the same core the live engine runs (no I/O), so dry-run and live route identically. Then exercise the **test harness** (`harness/`): its 5 headless `--scenario` runs

(`processed` / `filtered` / `unrouted` / `error` / `dead_letter`) assert dispositions over the API for CI, and its GUI can inject delivery faults (delay-then-AA, close, fail-N-then-AA) to prove your **retry / dead-letter / replay** behavior before you trust it.

Note: `validate` only catches **literal** port collisions; `env()` -resolved ports are checked at bind time. A `prod` -only missing `env()` value may not surface during a `dev` -context check — validate against the target environment before promoting.

I 9. Capacity & load testing on *your* hardware

Full references: [LOAD-TESTING.md](#) and the published [throughput baseline & tuning reference](#) (Gate #3) — a **two-tier gate**: host-independent **conformance** invariants (zero loss, bounded drain, low error rate — a hard release blocker) plus **performance** numbers "*as measured on the reference config*". Because the durable-write path is hardware-dependent, those msg/s figures are **not** a promise for your box — establish your own baseline.

The headless load harness (`harness/load/`) drives an already-running engine over real MLLP and the HTTP API (it never touches the store), so it is **store-agnostic** — swap the engine's `--db` to compare SQLite vs Postgres ceilings on identical traffic.

Recommended ramp:

1. **smoke** — tiny zero-loss wiring check (no performance claim).
2. **fanout-baseline** — warmup → ramp → sustained → spike → recovery; SLOs are evaluated only on the measured sustained phases. Reference targets in the profile: ≥200 msg/s sustained, ACK p99 ≤50ms, e2e p99 ≤5s, error ≤0.001, drain ≤60s, zero-loss.
3. **soak** — ~1-hour steady state watching DB/WAL growth + dead-letter accumulation.

Treat the **zero-loss reconciliation** (`sent == engine_read` , `sink_received == engine_written` , backlog drained to zero) as the **headline gate** — throughput numbers are meaningless if messages were lost. Use a **closed-loop** phase (fixed concurrency) to find your true max sustainable throughput, and the `slow` transform mode to find your per-core transform ceiling. Save the JSON/CSV reports and use `--baseline` + `--tolerance` to catch regressions over time. Size `correlator_capacity` above your peak in-flight (watch for correlation-miss notes), and remember a single Python sender process is the offer ceiling — shard it across processes if it can't saturate your engine.

Sizing reality: the staged pipeline has ~**3× write-amplification** on SQLite (3 commits for a common single-handler message; 2 + H for an H-way fan-out) — see [the write-amplification benchmark](#). Plan disk headroom for `.db` + `-wal` , plan retention/VACUUM (\$10), and move to **Postgres** if a single-writer SQLite ceiling becomes the bottleneck.

10. Backup, restore & disaster recovery

No existing repo doc covers this — it is part of *your* operational responsibility. Rehearse a full restore before you carry real data.

Back up the store.

- **SQLite:** the WAL backend means three files must be captured **consistently** — `.db`, `.db-wal`, `.db-shm`. Use `sqlite3 <db> ".backup '<dest>'"` against the live DB, or take a **quiesced cold copy** (graceful stop → copy → restart). A naive copy of just the `.db` while the service runs can be inconsistent.
- **PostgreSQL:** use your standard DB backups — `pg_dump` for logical backups and/or WAL archiving / PITR for point-in-time recovery. The engine is greenfield-only on server DBs, so the DB tier owns store-level DR here.

Escrow the encryption key SEPARATELY. If you enabled at-rest encryption (§6), a restored store is **unreadable without the same** `MEFOR_STORE_ENCRYPTION_KEY` / **DPAPI key file**. Back the key up in a different location/system from the data, with its own access control.

Restore-and-verify drill (do this in the lab, §11 Stage 0): restore the store + key into a clean host, start the engine, confirm `/health`, run `/status/integrity-check` (SQLite `PRAGMA quick_check`), and spot-check `/messages` and dispositions.

Keep the store bounded. `[retention]` is **off by default (kept forever)**. Set `max_db_mb` (drives a `storage_threshold` alert), `messages_days` / `dead_letter_days` (body purge), and the daily `VACUUM` so the store doesn't grow unbounded and a full disk doesn't take you down mid-pilot.

11. Staged rollout plan with go/no-go gates

This is the recommended path from first install to full production. **Do not skip stages** — each one exists to catch a different class of problem cheaply. Advance only when the stage's **exit criteria** are met.

Stage 0 — Lab / standalone

Goal: prove wiring, dispositions, and recovery on a throwaway box, with **synthetic data only**.

Setup: SQLite, loopback, auth on, a synthetic corpus from `messagefoundry generate`.

Exit criteria (→ Stage 1): - [] `messagefoundry check --config <dir> --messages <fixtures>` exits 0 (validate **and** dryrun green). - [] All 5 disposition `--scenario` runs pass (`processed` / `filtered` / `unrouted` / `error` / `dead_letter`). - [] You have driven a retry → dead-letter → **replay** cycle via the harness fault injection and understand the recovery tools (§7). - [] A **backup + restore** has been rehearsed once (§10).

Stage 1 — Shadow / parallel run

Goal: run MEFOR alongside your **incumbent** engine on **real production traffic** without affecting any downstream system.

How: tee/duplicate the production **inbound** feed to a MEFOR instance whose outbounds point at a **throwaway/null sink** (the harness correlation sink works well), or use a router that `Send`s only to a dedicated "shadow" outbound. Compare MEFOR's dispositions and transformed output against the incumbent's outcomes for the same messages.

 **Do not dual-write to real partners in shadow.** At-least-once + non-idempotent downstreams make a true dual-write dangerous. Keep shadow outbounds pointed at a sink unless the partner dedupes.

Exit criteria (→ Stage 2): - [] **Zero-loss reconciliation holds** over a sustained window (e.g. 1–2 weeks) at production volume. - [] MEFOR dispositions/output **match the incumbent's** for the same messages (differences explained). - [] **No unexplained dead-letters**; every `ERROR` understood. - [] A load test on **production-like hardware** meets your own SLO targets (§9).

Stage 2 — Limited production

Goal: MEFOR becomes the system of record for a **small, low-risk subset** of real feeds (one partner / one low-volume interface).

Prereqs: switch to **Postgres (single-node)** if you need a server DB; **encryption on** (`require_encryption=true`); **alerts wired** and **monitoring in place** (§12); **backups automated**; **upgrade + rollback runbook validated on staging** (§13).

Exit criteria (→ Stage 3): - [] e2e p99 within your SLO; **zero unexplained dead-letters** over the observation window. - [] **Alert wiring proven by a deliberate fault-injection drill** — you triggered `queue_buildup` / `connection_stopped` and the on-call was actually paged. - [] **Backup + restore rehearsed against the production store** (not just the lab copy). - [] Rollback runbook exercised at least once.

Stage 3 — Full production

Goal: migrate the remaining feeds **in waves** (never big-bang). Keep decommissioning the incumbent as a **separate, later** step so you retain a fallback.

Steady-state expectations: - [] Sustained-load SLO met on production hardware. - [] DR (backup/restore) rehearsed and scheduled. - [] On-call + the failure-drill runbook (§12) in place. - [] If you require HA: stand up the built **active-passive** cluster (leader/standby on a shared server-DB store — PostgreSQL or SQL Server) via the **§14 HA rollout runbook** — VIP standup + both failover drills passed — and keep HA at the DB tier too. Decided and rehearsed before you depend on it.

12. Day-2 operations & monitoring

Verify-it-runs (after every start/restart): `GET /health` → `{"status":"ok"}`, send a synthetic message, and confirm the **"wiring started"** banner in `service.out.log`.

Monitoring surfaces (poll the API + parse logs — there is no Prometheus exporter): - `/stats` — outbox counts by status. - `/status` — DB size vs disk free, journal mode, counts. **Scrape db-size-vs-disk-free.** - `/status/integrity-check` — on-demand store integrity. - `/ws/stats` — ~1 Hz queue-depth WebSocket. - `/messages` + `/messages/{id}` — per-message detail and the **event trail** (read this, not just the status — see §7). - `/dead-letters` — **page on dead-letter accumulation.** - **AlertSink events** to alert on: `connection_stopped`, `queue_buildup`, `storage_threshold` (wire these to webhook/email — §7). - `service.err.log` — watch it.

Note: the browser web console (`/ui`) surfaces dispositions, dead-letters, and alerts; the **CLI/API** remain available for scripted dead-letter triage and alert management.

Log management: logs land under `<DataDir>\logs` via NSSM. Configure rotation, keep the level at `INFO` or above (DEBUG can leak PHI — §6), treat `service.out/err.log` as **potential-PHI artifacts** (ACL them; don't ship them off-box — off-box logging is deferred), and include them in your retention policy.

Graceful drain for maintenance: stopping the service (Ctrl+C / NSSM stop) triggers the ASGI lifespan to call `engine.stop()` for a clean drain. Always **drain** → **stop** → **back up** → **change** → **restart** → **verify**.

Failure-drill runbook — rehearse these in the lab/staging before prod:

Symptom	First moves
Stuck FIFO lane (retry-forever head)	<code>queue_buildup</code> alert → inspect <code>/messages</code> for the head → fix-and- <code>replay</code> , or <code>dead_letter_now</code> to unblock the lane (the dead row stays replayable).
Poison message	It dead-letters under <code>CONTINUE</code> (or stops the lane under <code>STOP</code>) → triage via <code>/dead-letters</code> → fix the transform → replay.
Full disk / <code>storage_threshold</code>	Free space / tighten <code>[retention]</code> / <code>VACUUM</code> → confirm <code>/status</code> disk-free recovers.
Crash / unexpected restart	Startup auto-recovers in-flight rows → verify <code>/health</code> , the "wiring started" banner, and that backlog drains.
Planned maintenance	Drain-and-stop (graceful), perform the change, restart, verify.

13. Upgrade & rollback

Safe upgrade runbook (pinned-wheel model): 1. **Drain** inbound (quiesce senders or stop accepting new work) and confirm queues are draining. 2. **Stop** the service (graceful). 3. **Back up** the store **and** the encryption key (§10). 4. **Bump the pinned engine version:** update the pin in your config repo's `requirements.txt`

(`messagefoundry==<new>`) and `pip install "messagefoundry==<new>"` into the deployment venv. (A contributor on the editable install instead pulls the target commit/tag and `pip install -e .` — treat that checkout as the release artifact; §4.4.) 5. **Re-validate**: run `messagefoundry check` against your config (and `ruff / mypy / pytest` too if you develop the engine). 6. **Restart** and **verify** (`/health` , "wiring started" banner, `/status`).

Rollback: - **Config rollback** is the cheapest lever: the audited `POST /config/reload` does a quiesce-and-swap to a known-good `--config` directory (confined to the allow-listed reload roots). Keep your last known-good config dir available. - **Engine rollback**: re-pin the prior version (`pip install "messagefoundry==<prev>"`) → restart (same runbook above). (Contributors on the editable install: `git checkout` the prior commit/tag → `reinstall` → `restart`.) - ⚠ **Schema/store-level changes are not trivially reversible** against a populated store given the greenfield-only posture (no in-place migration). Plan code/config rollback as your primary path; use **dead-letter replay** to recover messages that a bad transform stranded before the rollback.

Pre-1.0 cadence: pin a released version (`messagefoundry==X.Y.Z`); the **latest release** is the supported target. Reproduce a problem against the latest release before filing an issue, and keep upgrades **small and frequent** rather than large and rare.

14. High availability — rollout & VIP standup runbook

Single-node is the default and is genuinely reliable — the durable staged queue (§7), not clustering, is what guarantees no message is lost on one node. Reach for HA only when you need **failover** (an unattended node loss must not stop intake), **not** for throughput (for that, scale intra-node — see the end of this section). When you do need it, MessageFoundry ships an **opt-in active-passive cluster** (Track B) — the supported HA model. This section is the operational runbook; the authoritative topology, lease semantics, and full settings catalog are in [CLUSTERING.md](#).

The model in one paragraph. Run N identical engine processes against **one shared server database** (PostgreSQL or SQL Server) with `[cluster].enabled = true`. Exactly one node — the **leader/primary** — runs the whole message graph (every listener *and* the router/transform/delivery workers); the rest are **warm standbys** that only contend for leadership (heartbeat + cache convergence), binding no listeners and running no workers until they win it. A **self-fencing leadership lease** in the shared DB elects the one primary and guarantees a partitioned old primary stops before a standby takes over. There is no separate broker and no node-to-node socket — coordination rides the shared `[store]` connection.

14.1 Stand up the cluster

1. **Pick a server-DB backend and make the DB tier itself HA**. `[store].backend = "postgres"` or `"sqlserver"` (SQLite is rejected for clustering). MEFOR coordinates the *processing* leader; it does **not** replicate your store — delegate store HA to the database (**PostgreSQL streaming replication / managed-Postgres HA**, or **SQL Server Always On**) and rehearse a DB failover separately.
2. **Provision every node identically**. Same installed engine version, the **same config dir**, and the **same** `[store]` **target** (server/database/schema) on each host. Set `[store].pool_size ≥ 2` (**≥ 3**

recommended — the leader drives the maintenance loop + reclaim sweep + workers concurrently).

3. **Enable the cluster** in the shared `messagefoundry.toml`: `toml [cluster] enabled = true`
`heartbeat_seconds = 10.0 # keep the invariant: heartbeat < fence < ttl`
`leader_fence_timeout_seconds = 20.0 # a leader that can't renew self-fences here (split-brain guard)`
`leader_lease_ttl_seconds = 30.0 # a standby may acquire only once the lease has expired`
The defaults trade a ~30 s crash-failover for margin; lower all three **proportionally** for faster failover at the cost of tolerance for a slow DB / GC pause.
4. **Sync clocks (NTP)**. Row leases use node wall-clock — keep nodes synced so a lease expiry isn't mistimed. (Leadership is evaluated on the *DB* clock, so it's skew-immune; the row leases are not.)
5. **Start the same `serve` command on every node** (run each as its own NSSM service — see [SERVICE.md](#)).
Confirm membership: `GET /cluster/nodes` lists every node and names exactly one `leader_node_id` ;
each node's `GET /cluster/status` reports `role: "primary"` on one and `"standby"` on the rest.

14.2 Stand up the VIP / load balancer (required)

Like Rhapsody/Corepoint, senders must reach "the engine" through a **floating VIP / L4 load balancer**, never a fixed node — because **only the primary binds the inbound listener ports**, the VIP is what makes a failover transparent (modulo a reconnect). **MEFOR ships the bind behavior and the `/cluster/*` endpoints, but does not ship a load balancer** — you stand one up (keepalived + IPVS, HAProxy in `tcp` mode, an F5, a cloud **L4 / network** LB, ...).

Data plane (MLLP / raw-TCP / X12 inbound) — one VIP per listener port:

- [] **Mode = L4 / TCP pass-through**. These are byte streams — do **not** front them with an HTTP/L7 LB.
- [] **Health check = a plain TCP connect to that exact listener port** on each backend node. Because only the primary binds the port, exactly one backend is ever healthy, so the VIP routes there automatically; on failover the old primary's port closes (goes unhealthy) and the new primary's opens (goes healthy) and the VIP follows. An application-level (MLLP-message) probe is unnecessary — a TCP connect is the correct, sufficient signal.
- [] **Tune the probe to your failover budget**. The VIP repoints in roughly `check_interval × unhealthy_threshold` ; size it well under your acceptable outage and above your DB/GC jitter (a 2–5 s interval is typical).
- [] **If MLLP-over-TLS is on, keep TLS end-to-end** (TCP pass-through) so the engine's own cert is presented; don't terminate TLS at the VIP unless you re-encrypt to the backends.
- [] **Don't let the LB reap long-lived MLLP sockets** — set TCP idle timeouts generous enough for persistent connections.
- [] **Verify every partner reconnects-on-drop**. On failover, connections to the old primary drop and senders must redial the VIP. This is standard MLLP client behavior — confirm it per partner.

Management plane (console / IDE → engine API) — optional VIP:

- [] The API is a control/read plane over the shared DB and is **up on every node**, so an API VIP can health-check the unauthenticated `GET /health` (liveness); you can point the console at any node.

- [] To pin operators to the active leader, read `GET /cluster/status` → `role` (`primary` / `standby`) or `GET /cluster/nodes` → `leader_node_id` / `lease_owner` (the console already surfaces the live primary). Both endpoints need only `MONITORING_READ` (VIEWER and up; no PHI).

14.3 Rehearse & validate failover (before real traffic)

Failover is **not instantaneous** — quantify *your* window from these drills; don't assume zero-downtime.

- [] **Clean switchover** (planned): gracefully stop the primary's service. It **expires its lease**, so a standby promotes on its next heartbeat ($\approx$ one `heartbeat_seconds`). Watch `leader_node_id` move on `/cluster/nodes`, watch the VIP repoint, and keep synthetic traffic flowing throughout.
- [] **Crash** (unplanned): hard-kill / power off the primary. Its lease **ages out**, so a standby promotes after **up to** `leader_lease_ttl_seconds` (~30 s default); a partitioned old primary **self-fences within** `leader_fence_timeout_seconds` so it can't double-process. Confirm the new primary's **owner-scoped on-promotion recovery** re-pends the dead primary's in-flight rows immediately (delivery resumes without waiting out the per-row lease TTL).
- [] **Zero-loss across the failover**. At-least-once means a row interrupted mid-delivery is **re-delivered** after its lease expires — so confirm your downstream connectors **dedupe / are idempotent**, then reconcile sent-vs-delivered across the drill.
- [] **Record the measured window** and, if needed, tune `heartbeat < fence < ttl` down proportionally and re-drill.

14.4 HA go-live checklist

- [] Server-DB backend; **DB-tier HA configured and its own failover rehearsed**.
- [] Identical engine version + config dir + `[store]` target on every node; `pool_size ≥ 3`; **clocks NTP-synced**.
- [] `/cluster/nodes` shows all nodes and exactly one `leader_node_id`.
- [] **VIP per inbound port** with a **TCP-connect health check**; partner **reconnect-on-drop** verified.
- [] **Clean-switchover drill** passed ($\approx$ one heartbeat, VIP follows, zero loss).
- [] **Crash drill** passed (failover within the TTL, no double-processing, in-flight rows recovered, zero loss).
- [] Measured failover window documented; lease timings tuned to your network.
- [] **Config changes applied as a coordinated (non-rolling) restart** — all nodes restart together so they never run divergent graphs across the change window.
- [] `/cluster/status` + `/cluster/nodes` monitored (alert on **no live leader** beyond a failover window) alongside the §12 surfaces.

For throughput (not failover), scale intra-node: one independent delivery worker per outbound connection (a slow/failing lane never blocks siblings), and keep retry policies finite where head-of-line blocking on a shared FIFO lane would otherwise stall throughput. Adding cluster nodes does **not** raise throughput — only the leader processes.

| 15. Getting help & reporting bugs

- **Bugs & feature requests:** open a GitHub issue using the repository's issue templates (`bug_report.md` / `feature_request.md`).
 - **Security vulnerabilities:** use the repository's **private security advisory** process per `.github/SECURITY.md` — do **not** open a public issue for a vulnerability.
 - **Before filing:** verify against current `main` (pre-1.0, latest-main-only support), and include the engine version, config shape, and relevant **non-PHI** log excerpts.
 -  **Never attach real PHI** to an issue, log excerpt, or reproduction. Reproduce with a synthetic corpus from `messagefoundry generate` .
-

| 16. Decommissioning a pilot

Ending a pilot is a **PHI-disposal** event. `uninstall-service.ps1` removes the service but **deliberately leaves the store and logs on disk**. To tear down cleanly:

1. **Graceful drain + stop**, confirm no in-flight work remains.
2. **Uninstall the service** (`scripts\service\uninstall-service.ps1`).
3. **Securely dispose of all PHI-bearing artifacts:** the store (`.db` + `-wal` + `-shm`), any PostgreSQL database/backups, **File-connector spill directories**, the `logs` directory, every **backup copy**, and the **encryption key / DPAPI key file**.
4. **Revoke credentials** (service account, AD bind account, any API tokens).

Treat backups and the encryption key with the same disposal rigor as the live store — a forgotten encrypted backup plus its escrowed key is still recoverable PHI.

Appendix — where each topic is documented

Topic	Reference
Install + your config repo (consumer model)	INSTALL-GUIDE.md , ADR 0017
System requirements / sizing by volume	SYSTEM-REQUIREMENTS.md
Install + Windows service	SERVICE.md
Network exposure / TLS	DEPLOYMENT.md
High availability / clustering	CLUSTERING.md
Throughput baseline / tuning	TUNING-BASELINE.md
Service settings / environments	CONFIGURATION.md
Connections / the graph / <code>connections.toml</code>	CONNECTIONS.md , ADR 0007
Reliability / staged pipeline	ADR 0001
Security / auth / RBAC / audit	SECURITY.md
PHI handling / encryption	PHI.md
HL7 validation	HL7-VALIDATION.md
Load testing	LOAD-TESTING.md
Write-amplification / sizing	step-b-write-amplification.md
Built-vs-roadmap (authoritative)	ARCHITECTURE.md