

System Requirements & Sizing

These are the minimum and recommended requirements for running the MessageFoundry engine, its message store, and the administration clients. The engine is a headless Python/asyncio service; the console is a separate desktop application.

On the throughput figures. MessageFoundry does not publish a validated, guaranteed per-node throughput baseline. The sizing tiers in [Sizing by message volume](#) are **engineering estimates** derived from the architecture and from synthetic load-test profiles — they are starting points, not guarantees. Always establish your own baseline on production-like hardware before go-live (see [Capacity notes](#)).

Hardware

	Minimum (lab / low-volume pilot)	Recommended (single-node production)
CPU	2 cores	4+ cores (transform throughput is per-core; one worker set per connection)
Memory	4 GB	8–16 GB
Disk	10 GB free, any disk	SSD, 50+ GB or sized to your retention window, on a low-latency local volume
Store volume	—	Put the message store on a fast local disk (not a network share). Budget for store + WAL growth : the staged pipeline writes roughly 3× per message on the embedded store.

The engine and store may share a host for a low-volume pilot. For production, run a **server database** (PostgreSQL or SQL Server) on its own host, sized by your DBA, and keep the engine host dedicated. For volume beyond one CPU core, see [Sizing by message volume](#).

Operating systems

Platform	Status
Windows Server 2022 / 2025	Primary supported and serviced platform (Windows-service deployment)
Windows Server 2019	Supported
Windows 10 / 11	Supported (development, pilot, console host)
Linux (modern x86-64 distributions)	Engine supported (cross-platform Python); no bundled service installer — run under systemd yourself
macOS	Development / console use only

Runtime

Component	Requirement
Python	3.14 or later , 64-bit (3.14 supported and CI-validated)
Service manager (Windows)	A Windows service wrapper, auto-provisioned and SHA-256-pinned by the installer, or pre-staged. Requires administrator / elevation to register the service.
C compiler	Not required for the default install (runtime dependencies ship as wheels)

Databases (message store)

Database	Status	Driver / prerequisite
SQLite (WAL)	Default, bundled — single-node	None (<code>aiosqlite</code> , in-process)
PostgreSQL 13+	Production	<code>messagefoundry[postgres]</code> extra (<code>asyncpg</code> , pure-Python — no OS dependency)
Microsoft SQL Server 2019 / 2022	Production	<code>messagefoundry[sqlserver]</code> extra (<code>aiodbc</code>) plus the OS-level Microsoft ODBC Driver 18 for SQL Server . Read-Committed Snapshot Isolation (RCSI) recommended.
MySQL / Oracle	Not supported	—

The embedded SQLite store needs no setup and suits pilots and single-node deployments. A **server database is greenfield-only** — there is no in-place migration from a populated SQLite store; drain and cut over. The database tier owns its own backup, HA, and (SQL Server) TDE / purge maintenance. A server database is also the **concurrency / scale substrate** — see below.

Administration clients

Client	Requirement
Desktop console	PySide6 (Qt) desktop application — install with the <code>console</code> extra. Runs on Windows, Linux, or macOS as a separate process; connects to the engine over the HTTP/WebSocket API. Not browser-based.
VS Code extension	Visual Studio Code (current stable) — route wizard, validate-on-save, test bench, stage → promote.
Web browser	Not required to operate the engine (no web admin UI; the API is HTTP/JSON for tooling).

Network & ports

Purpose	Default	Notes
Engine API (HTTP + WebSocket)	127.0.0.1:8765	Loopback by default , authentication required. To reach the API from another host, front it with a TLS-terminating reverse proxy or tunnel — exposing the listener off loopback should always be done over TLS.
Inbound MLLP / TCP listeners	operator-defined (samples use e.g. 2575 , 2600)	Open to sending systems via firewall. Keep MLLP on a trusted network segment .
Outbound	as configured	Reachability to downstream partners and, for server databases, to the database host.
Installer egress	HTTPS	Outbound access for the service installer to fetch the pinned service-wrapper binary (or pre-stage it).

Interface authentication. Beyond network-segment controls, interfaces can authenticate with **mutual TLS (mTLS)**, **OAuth 2.0 client-credentials**, and **SMART-on-FHIR Backend Services**. Administrative access uses required multi-factor authentication for local accounts; enterprise / Active Directory users authenticate (and get MFA) through their own identity provider via SSO federation.

Sizing by message volume

Engineering estimate — not a validated benchmark. These tiers project the architecture's behavior; they are not committed numbers. Throughput depends heavily on **transform cost per message** (the dominant factor), message size, fan-out, and strict-validation use. **Measure your own feeds** with the load harness before committing (see [Capacity notes](#)).

How throughput is bounded (read this first)

A single engine process runs **all** message work — decode → peek → route → transform → re-encode — on **one CPU core** (one asyncio event loop; the GIL prevents pure-Python parallelism across threads). So per-process throughput is governed, in order, by:

1. **Transform cost per message** — usually the binding constraint. As an illustrative estimate, a comparable vendor benchmark shows real transformation cutting pass-through throughput by roughly 60% (≈ 1000 msg/s → ≈ 400 msg/s). A light / pass-through feed sits near the top of a tier; a heavy transform sits near the bottom.
2. **Durable-write cost** — every stage handoff (ingress → routed → outbound → delivered) is a committed transaction. In-process **SQLite is fastest per write**; a **server database is slower per single write** (network + MVCC) but is the concurrency substrate (next point).

To exceed one core, scale **intra-node** on a server database: many connections / lanes / delivery workers draining **one shared server database** (PostgreSQL or SQL Server) concurrently via `SELECT ... FOR UPDATE SKIP LOCKED` plus row leases. Throughput scales with workers until the **database's commit capacity** is the wall. (SQLite is single-writer and does **not** scale this way — it is the single-process / single-node store.)

Engine HA is **single-leader active-passive** — the graph runs on the leader only. A **multi-process, sharded-by-inbound** scale-out (multiple engines, each owning a disjoint set of inbounds, on the shared database) is a **future direction, not a current capability**.

Tiers

Tier	Peak sustained (est.)	Indicative daily volume	Deployment shape	Store	Suggested hardware (engine host)
Pilot / light	up to ~50 msg/s	up to ~1–4 M/day	1 process, single node	SQLite	2 cores / 4 GB
Standard single-node	~50–200 msg/s	~4–15 M/day	1 process, single node	SQLite, or PostgreSQL / SQL Server	4 cores / 8 GB
High single-node	~200–500 msg/s	~15–40 M/day	1 process, tuned (lean transforms; finite-retry on hot lanes)	Server database recommended (PostgreSQL / SQL Server)	4–8 cores / 16 GB
High single-node, concurrent	up to ~500 – low-thousands msg/s	up to ~40 M+/day	1 process, many connections / lanes draining concurrently via <code>SKIP LOCKED</code>	PostgreSQL / SQL Server (required — not SQLite)	8+ cores / 32 GB + a dedicated database host sized to the commit load

Reading the tiers

- *Peak sustained* is a **per-second** capacity estimate. Healthcare feeds are bursty; real **average** rate (and therefore daily volume) is typically a fraction of peak, so the *indicative daily volume* columns assume a realistic duty cycle, not $\text{peak} \times 86,400$.
- The **single-stream / single-core ceiling** is roughly the "High single-node" row — a few hundred msg/s with real transforms, approaching ~1000 msg/s only for light / pass-through work. Past that on one feed you are over one core's budget.
- The **estimated maximum as currently architected** is the "**High single-node, concurrent**" row: one engine process, many connections / lanes draining the shared server database concurrently via `SKIP LOCKED`, bounded by the database's commit ceiling. There is no fixed published cap — on a well-provisioned box with a tuned server database this lands in the **low thousands of msg/s**, beyond which you are **database-bound** and scale the database tier. (A multi-process, sharded-by-inbound scale-out beyond one engine is a **future direction**.) Group-commit and a lazy MSH-only routing peek are identified levers to raise the per-core ceiling.

Single-stream server-DB caveat. Because each staged handoff is a committed round-trip, a single delivery worker against a *remote* server database drains far slower than in-process SQLite (a SQL Server CI smoke profile observes roughly 30 deliveries/s for one stream). High volume on a server database comes from **concurrency** — many connections / lanes /

processes draining in parallel — not single-stream speed. Size the database host for that concurrent commit load.

Capacity notes

- Validate with the load harness: run a `smoke` → `fanout-baseline` → `soak` ramp, exercise the `cheap` / `edit` / `slow` transform modes to find your per-core transform ceiling, and compare SQLite vs a server database on identical traffic. Treat the **zero-loss reconciliation** as the headline gate — throughput is meaningless if messages were lost.
- The embedded store has roughly 3× write amplification and a single-writer ceiling; move to PostgreSQL or SQL Server when that becomes the bottleneck.
- Scale **intra-node** on a server database (one delivery worker per outbound; many connections / lanes draining concurrently via `SKIP LOCKED`; keep retry policies finite where head-of-line blocking on a shared FIFO lane would otherwise stall a lane). A multi-process scale-out beyond one engine is a **future direction**. Engine **HA** is **active-passive failover** (opt-in leader / standby cluster on shared PostgreSQL — see [CLUSTERING.md](#)); delegate **database-tier** HA to the database plus a load-balancer VIP.

MessageFoundry is independent and unaffiliated with Mirth, Corepoint, Cloverleaf, Rhapsody, or Ensemble; those names are trademarks of their respective owners.