

---

# MessageFoundry User Guide

**Early Access** · MessageFoundry v0.3.2 · July 2026

MessageFoundry is an open-source, Python interface engine for healthcare — a modern alternative to Mirth Connect and Corepoint. It **receives, routes, transforms, and validates** messages between systems (HL7 v2.x by default, payload-agnostic for other formats), with routing and handling expressed as **ordinary Python you own and version-control**. This guide is **task-oriented** ("how do I..."): it walks you from a clean machine to a running engine, then through authoring Connections/Routers/Handlers and operating and monitoring the system. It links to the reference docs rather than restating them.

## Contents

---

- What MessageFoundry is, and how to use this guide
- Installing and running the engine
- Quickstart: send your first message
- Authoring Connections
- Authoring Routers and Handlers
- Operating with the console and the VS Code extension
- Monitoring dispositions and troubleshooting
- Where to go next

---

## What MessageFoundry is, and how to use this guide

MessageFoundry receives, routes, transforms, and validates messages between systems: **HL7 v2.x by default**, and payload-agnostic for other formats (JSON, XML/SOAP, X12 EDI, database records). Unlike a legacy engine's embedded scripting language or a locked-in low/no-code GUI, the routing and handling logic is **ordinary Python you own and version-control** — and connection setup in particular can be pure data (a TOML file, edited by hand or in a VS Code GUI). The engine runs headless as an asyncio service; you operate it from a browser **web console** (served same-origin at `/ui`) and a VS Code extension, both talking to it over a localhost HTTP/WebSocket API.

### The mental model in one read

The configuration is a **graph wired by name** — there is **no "channel" object** that bundles a source, filters, transforms, and destinations the way a legacy engine does. Instead you wire four building blocks, using the project's exact vocabulary:

- **Connection** — an endpoint that **receives** ( `inbound()` ) or **sends** ( `outbound()` ) messages: MLLP, TCP, File, REST, SOAP, Database, SFTP/FTP. Every message in or out is counted and logged.
- **Router** ( `@router` ) — a pure Python function bound to **one** inbound Connection. It sees every received message and returns the **name(s)** of the Handler(s) to forward to (or none, to filter).
- **Handler** ( `@handler` ) — a pure Python function that takes a message, **filters** → **transforms**, then returns `Send(...)` s naming one or more outbound Connections.
- **Message store** — durable persistence + the staged queue. SQLite (WAL) by default; PostgreSQL or SQL Server for production.

The edges between these are just names resolved at config load: an inbound names its Router, a Router returns Handler names, a Handler `Send` s to a named outbound. To understand a feed, follow the names; to change it, change a name. (See `samples/config/IB_ACME_ADT.py` and `samples/config/adt.py` for a complete, runnable example of this wiring.)

Under the hood, a received message flows through a **staged pipeline** of three persisted stages — **ingress** (the raw message, committed *before* the ACK) → **routed** (one row per Handler the Router selected) → **outbound** (one row per destination) — each drained by its own asyncio worker. Because nothing is ever silently dropped, every message carries a **disposition** that the store records as it flows: `RECEIVED` → `ROUTED` / `UNROUTED` → `PROCESSED` / `FILTERED` / `NOT_DEPLOYED` / `ERROR` . This is the count-and-log promise; operators read the disposition (and alerts), not the ACK, to confirm a message reached its destination. The full reliability model (at-least-once delivery, why routers/handlers must be pure) is in [ADR 0001](#).

## Who this guide is for

This guide serves two audiences: **operators** who install, run, and monitor the engine, and **config authors** (integration developers and analysts) who wire Connections, Routers, and Handlers. It links out to the reference docs rather than restating them. If you want a lighter, narrative onboarding first, read [EARLY-ADOPTER-GUIDE.md](#) (install-to-production rollout) or [MENTAL-MODEL.md](#) (the concepts above, in depth).

## Where the depth lives — the reference map:

| When you want...                                                                 | Read                                                                                 |
|----------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| The full architecture, modularity standard, and dependency rules                 | <a href="#">ARCHITECTURE.md</a> (diagrams: <a href="#">architecture-diagram.md</a> ) |
| Connection types, settings, and the graph (incl. <code>connections.toml</code> ) | <a href="#">CONNECTIONS.md</a> , <a href="#">ADR 0007</a>                            |
| Translation tables (code sets) — the grid editor + <code>codeset</code> CLI      | <a href="#">CODESETS.md</a> , <a href="#">ADR 0033</a>                               |
| Service settings, environments, and <code>env()</code> values                    | <a href="#">CONFIGURATION.md</a>                                                     |
| Running as a Windows service (NSSM)                                              | <a href="#">SERVICE.md</a>                                                           |
| Auth, RBAC, audit, and TLS                                                       | <a href="#">SECURITY.md</a> , <a href="#">DEPLOYMENT.md</a>                          |
| PHI handling and encryption-at-rest                                              | <a href="#">PHI.md</a>                                                               |

| When you want...                                                                                        | Read                                   |
|---------------------------------------------------------------------------------------------------------|----------------------------------------|
| The staged pipeline / reliability; payload-agnostic ingress; the read-only <code>db_lookup</code> ; X12 | ADR 0001, ADR 0004, ADR 0010, ADR 0012 |
| What's built vs. planned                                                                                | FEATURE-MAP.md, README.md              |

A note on PHI before you run anything: this engine carries PHI, and a few commands (notably `dryrun` and `generate`) print **full message bodies** to stdout. Use **synthetic HL7 only** in examples, and never redirect that output to a committed file, ticket, or CI log. To build *realistic* test data from real traffic without ever handling PHI, use the built de-identification framework via the standalone tee relay's `python -m tee anonymize-captures` (fail-closed; ADR 0030) rather than hand-writing synthetic messages. See [PHI.md](#) for the hard rules.

## How the rest of this guide is laid out

The remaining sections are a **path**, in the order you'll actually work through them:

1. **Install** the engine and scaffold your config repo.
2. Send your **first message** end to end (e.g. `samples/messages/adt_a01.hl7` via `samples/send_mllp.py`) against the bundled sample config.
3. **Author Connections** (in Python or `connections.toml`).
4. **Author Routers and Handlers** to route and transform.
5. **Operate** via the browser web console ( `/ui` , `messagefoundry-webconsole`) and the VS Code extension (`ide/`).
6. **Monitor and troubleshoot** with dispositions, alerts, dead-letter triage, and replay.

Work through them in order the first time; afterward, jump to the task you need.

## Installing and running the engine

This section takes you from a clean machine to a running engine and an attached console. It covers the **developer / checkout** path (running the bundled `samples/config`) and points at the **pinned-wheel** consumer path where they differ.

Two install styles exist. To **try MessageFoundry against the sample config** (this guide's running example), install from a checkout ( `pip install -e .` ). To **deploy your own interfaces**, install the pinned wheel and scaffold a config repo with `messagefoundry init` — the full consumer model is [INSTALL-GUIDE.md](#).

### 1. Check prerequisites

- **Python 3.14+** (64-bit). Everything else (the SQLite store, NSSM for the service) is auto-provisioned or bundled.
- **git** if you are cloning the checkout or standing up a config repo.
- **Administrator/elevation** only for the Windows service step (6).

Full hardware, OS, database, and port details: [SYSTEM-REQUIREMENTS.md](#).

## 2. Install the engine

Create a virtual environment and install. For the **checkout / contributor** path (gives you `samples/config`, `samples/messages/`, and `samples/send_mllp.py`):

```
python -m venv .venv
.\.venv\Scripts\Activate.ps1      # Linux/macOS: . .venv/bin/activate
pip install -e .                  # core runtime + SQLite store
```

For a **deployment**, pin the published wheel instead (and verify its provenance — see [INSTALL-GUIDE.md](#)):

```
pip install "messagefoundry==0.3.2"
```

Add only the extras a host actually needs (each is opt-in and lazy-imported):

```
pip install -e packaging/messagefoundry-webconsole # the browser web console (/ui) – the operator UI (step 5)
pip install -e ".[harness]"                        # the standalone PySide6 test harness GUI
pip install -e ".[postgres]"                       # PostgreSQL store backend
pip install -e ".[sqlserver]"                      # SQL Server store backend – also needs the OS-level Microsoft ODBC Driver 18
pip install -e ".[sftp]"                           # SFTP transport for the REMOTEFILER connector
pip install -e ".[fhir]"                           # FHIR codec (R5/R4B/STU3) + FHIR() outbound (ADR 0022)
pip install -e ".[dicom]"                          # DICOM C-STORE SCP + codec – headers/SR only, no pixels (ADR 0025)
pip install -e ".[webauthn]"                       # browser WebAuthn passkeys for the /ui console (ADR 0068)
pip install -e ".[otel]"                          # OpenTelemetry/OTLP export seam (the /metrics endpoint itself needs no extra)
```

(For a deployment wheel, the same extras apply: `pip install "messagefoundry[harness]==0.3.2"`, and the web console installs as its own wheel `pip install "messagefoundry-webconsole==0.2.15"`.) SQLite is the zero-dependency default — you need no extra to run the sample config.

## 3. Run the engine headless (dev)

From the repo root, against the bundled sample config:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

- `--config` points at the directory of Connection/Router/Handler modules (here `samples/config/`, which includes `IB_ACME_ADT.py` and its transform `adt.py`).
- `--db` is the SQLite message store path (created on first run).
- `--env` is required — `serve` refuses to start without it. The active environment is a free-form **name** (`dev` / `staging` / `prod`, or a custom name) that does two things: it selects the value file `environments/<env>.toml` that `env("...")` lookups resolve against, and it sets the instance's **PHI posture**

( `data_class` / `production` ). Built-in names carry a default posture; a custom name must declare it. See [CONFIGURATION.md](#).

When the engine runs from somewhere other than the repo root (e.g. under the service), anchor the value files with `--project-root <repo-root>` so `env()` values don't silently resolve empty — see [INSTALL-GUIDE.md](#).

**Network / auth posture.** The API binds `127.0.0.1:8765` and **requires authentication** by default. A non-loopback bind without TLS is refused at startup; configure native TLS (or an upstream terminator) to expose it. Details: [SECURITY.md](#) and [DEPLOYMENT.md](#).

**Store encryption (PHI instances).** On a PHI instance ( `[security].handles_real_patient_data = true` ), `serve` **refuses to start** — in every environment, `dev` and `staging` included, not just production — if no store encryption key is configured. Mint one with `messagefoundry gen-key` (set it as `MEFOR_STORE_ENCRYPTION_KEY` ), or on Windows DPAPI-protect it to a file with `messagefoundry protect-key --generate --out <file>` and point `[store].encryption_key_file` at it. The full key story is in [PHI.md](#).

Confirm it's up:

```
curl http://127.0.0.1:8765/health
```

## 4. Scaffold your own config repo

When you're ready to author real interfaces, scaffold a standalone, `check` -green config repo instead of editing the samples:

```
messagefoundry init ./my-config-repo
cd my-config-repo
```

It writes a runnable starter feed under `config/`, `environments/dev.toml` + `prod.toml`, a synthetic fixture under `messages/sets/`, a `messagefoundry.toml`, a `requirements.txt` pinning the engine, and a CI `check.yml`. Validate and run it:

```
pip install -r requirements.txt
messagefoundry check --config config --messages messages/sets
messagefoundry serve --config config --env dev
```

Put it under version control with **Set Up Version Control & Checks** in the IDE (or a plain `git init`). During setup you choose **where the repo is stored** — *on this machine only* (fine for a single-box, non-HA engine or local dev) or on a *shared remote* (recommended for HA, a team, or off-machine backup); change it any time with **Config Repo Storage Location**. Choosing local vs. remote storage, and why secrets/PHI never land in the repo, are covered in [VERSION-CONTROL.md](#) (with the full deployment model in [INSTALL-GUIDE.md](#)).

## 5. Open the admin console (in a browser)

The console is the **browser web console** served same-origin by the engine at `/ui` — install the `messagefoundry-webconsole` wheel alongside the engine and it is served automatically, **on by default** at a loopback bind (turn it off with `[security].serve_web_console = false`). With the engine running, browse to:

```
http://127.0.0.1:8765/ui
```

The web console prompts for sign-in (authentication is on by default). Source: `packaging/messagefoundry-webconsole/`. (The former PySide6 desktop console was retired — BACKLOG #103; PySide6 now backs only the standalone test harness.)

## 6. Run as a Windows service (NSSM)

For production on Windows, run the engine as a background service via **NSSM** — it starts on boot, restarts on crash, captures stdout/stderr to rotating logs, and stops with Ctrl+C so connections drain cleanly. From an **elevated** PowerShell:

```
.\scripts\service\install-service.ps1 -Environment prod
```

`-Environment` is **required** (it becomes `serve --env`, just like step 3). The script is idempotent (re-run to reconfigure), auto-downloads a SHA-256-pinned NSSM if one isn't on `PATH`, and defaults to service name `MessageFoundry`, config `<repo>\samples\config`, store + logs under `C:\ProgramData\MessageFoundry`, bind `127.0.0.1:8765`. Override paths/port/account with flags, e.g.:

```
.\scripts\service\install-service.ps1 -Environment prod -Port 9000 `
  -Config D:\hl7\config -DataDir D:\MessageFoundry `
  -ServiceAccount "NT SERVICE\MessageFoundry" # least-privilege; auto-grants the needed ACLs
```

Manage and remove it:

```
nssm start MessageFoundry
nssm status MessageFoundry
nssm stop MessageFoundry
.\scripts\service\uninstall-service.ps1 # elevated; leaves logs + store in place
```

The complete procedure — least-privilege accounts, locking down the config/log directories, DPAPI key protection, update-vs-reinstall, and troubleshooting — is in `SERVICE.md`.

## A note on PHI-emitting commands

`messagefoundry dryrun` and `messagefoundry generate` print **full message bodies** to stdout/stderr (`dryrun` only with `--show-phi`; redacted otherwise). Run them against **synthetic HL7 only** — never real PHI — and never redirect their output into a committed file, ticket, or CI log. See `PHI.md`.

## Quickstart: send your first message

This walkthrough uses only the shipped samples in `samples/config/` — no editing required. You'll start the engine, push a synthetic HL7 ADT message over MLLP, and watch it get received, routed, and archived to a file.

The sample inbound that does the work is `IB_Test_ADT` in `samples/config/adt.py`: an MLLP listener on **port 2575** whose Router forwards `ADT` messages to the `archive` Handler, which writes them to `./out/adt/{MSH-10}.hl7` via the `FILE-OUT_Test_ADT` outbound connection. (The config dir also wires other sample feeds — ACME ADT on 2600, X12, immunizations — but this Quickstart only exercises `IB_Test_ADT`.)

### 1. Start the engine

In one terminal, from the repo root, run the engine against the sample config. The active `--env` is **required** — use `dev` for local work:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

This loads the config modules, opens (or creates) the SQLite store at `./messagefoundry.db`, serves the localhost API, and starts every sample listener. The startup log announces the active environment and posture. Leave it running. The API binds `127.0.0.1` by default; first-run auth/console details are covered in the install guide — see `INSTALL-GUIDE.md`.

### 2. Send the sample message

In a **second** terminal (also from the repo root), send the shipped synthetic ADT^A01 over MLLP with the helper in `samples/send_mllp.py`:

```
python samples/send_mllp.py samples/messages/adt_a01.hl7
```

The script defaults to `--host 127.0.0.1 --port 2575`, which matches `IB_Test_ADT`, so no flags are needed. To be explicit (or to target a different listener), pass them:

```
python samples/send_mllp.py samples/messages/adt_a01.hl7 --host 127.0.0.1 --port 2575
```

The file it sends, `samples/messages/adt_a01.hl7`, is synthetic — never substitute real PHI here.

### 3. What to expect

`send_mllp.py` reuses the engine's MLLP framing, waits for the framed ACK, and prints it:

```
--- ACK ---
MSH|^~\&|...|...|...|ACK^A01|...|P|2.5.1
MSA|AA|MSG00001
...
```

An `MSA|AA` is a positive acknowledgement ( `AA` = Application Accept). Under the staged pipeline the ACK means **received and durably persisted** (status `RECEIVED` at the ingress stage), *not* final delivery — see the ACK-on-receipt model in `ARCHITECTURE.md`. A moment later the message is routed and delivered: because `adt_a01.h17` is an `ADT^A01`, the Router forwards it to the `archive` Handler, the Handler stamps a downstream facility mnemonic *if one is mapped for the sending facility* (none is for this sample, so its MSH-4 is left unchanged) and sends it on, and the file outbound writes it under `./out/adt/`. Its disposition advances `RECEIVED` → `ROUTED` → `PROCESSED`. (A non-`ADT` message would be logged `UNROUTED`; an `ADT` event the Handler drops would be `FILTERED` — nothing is silently discarded.)

Confirm the delivered file landed (its name is the message's MSH-10 control ID, `MSG00001`):

```
ls ./out/adt/
```

## 4. Where to see it land

Two complementary ways to inspect what happened:

- **The web console Messages page.** Open the web console at `/ui` in a browser, then use **Traffic** → **Messages** to find the message, its disposition, the original raw body, and the per-destination delivery. The web console talks only to the engine's API — see `packaging/messagefoundry-webconsole/`.
- **A dryrun preview (no engine needed).** To see exactly how the config *would* route and transform a message without sending it anywhere, run:

```
python -m messagefoundry dryrun --config samples/config --messages samples/messages/adt_a01.h17 --inbound IB_Test_ADT
```

Dryrun prints the resolved inbound, disposition, selected handlers, and would-send payloads ( `--inbound` names which inbound to evaluate — required here because the sample config wires several). **Its output is PHI-bearing** — message bodies are redacted by default and only included with `--show-phi`. The samples here are synthetic, but never pipe `dryrun` (or `generate`) output to a committed file or a CI log.

## Now change what it does

Once the round trip works, make it yours: edit the Router and Handler in `samples/config/adt.py` to change routing and transforms (see `Authoring Routers and Handlers`), add or retarget connections in `samples/config/connections.toml` — by hand or via the VS Code editor (`ADR 0007`) — and follow the connection naming convention in `CONNECTIONS.md`. Run `python -m messagefoundry check --config samples/config --messages samples/messages` before committing config changes.

---

## Authoring Connections

A **Connection** is an endpoint that either *receives* messages (an **inbound** source) or *sends* them (an **outbound** destination). MLLP and File are the two most common transports, both shipped today (plus raw TCP, X12, REST, SOAP, Database, and SFTP/FTP — see the full catalog and per-setting reference in `CONNECTIONS.md`).



Connections carry only *transport* config; routing/transform *logic* lives in code-first Routers and Handlers (see [Authoring Routers and Handlers](#)).

You author a connection one of two ways — **as code** (a `.py` module) or **as data** ( `connections.toml` ). Both desugar into the same registry, so they coexist freely.

## Name your connection

Use the convention `[TYPE]_[PARTNER]_[MESSAGE]` :

- **TYPE** — transport + direction code: `IB` / `OB` (inbound/outbound MLLP), `FILE-IN` / `FILE-OUT` , `TCP-IN` / `TCP-OUT` , `X12-IN` / `X12-OUT` , etc. (the full table is in `CONNECTIONS.md`).
- **PARTNER** — the system on the other end ( `ACME` , `Epic` , `Test` ).
- **MESSAGE** — the HL7 message code ( `ADT` , `ORU` , `VXU` , ...) or `MIXED` / `ALL` .

Example: `IB_ACME_ADT` = inbound MLLP from ACME carrying ADT. Names are plain strings, so hyphens and mixed case ( `FILE-OUT_Test_ADT` ) are fine. Router/Handler names are *not* connections and don't follow this formula.

## Author a code-first inbound and outbound

In a config module, call the `inbound()` / `outbound()` factories with a transport spec ( `MLLP()` , `File()` , ...). Here is the shape, drawn from `samples/config/IB_ACME_ADT.py`:

```
from messagefoundry import MLLP, env, inbound, outbound

inbound("IB_ACME_ADT", MLLP(port=2600), router="acme_adt_router")
outbound("OB_ACME_ADT", MLLP(host=env("acme_adt_host"), port=env("acme_adt_port", cast=int)))
```

Key points the sample demonstrates:

- **Inbound MLLP takes only a port** — passing `host` is a wiring error. The listen interface is the service-level `[inbound].bind_host` (loopback in DEV, a specific NIC in PROD), an operator setting, not authored here.
- **Outbound MLLP needs a host and port** . Anything that differs by environment (a downstream peer, a credential) uses `env("key")` , resolved per instance from `environments/<env>.toml` (and `MEFOR_VALUE_<KEY>` for secrets) — so one module runs unchanged in every environment. A referenced-but-undefined value fails loud at load, never a silent blank host.
- For a **File** endpoint, use `File(directory="./out/adt")` (in) / `File(directory=..., filename="{MSH-10}.h17")` (out). For non-HL7 bodies, set the inbound's `content_type` so the body routes as a `RawMessage` instead of being HL7-parsed — the shipped set is `hl7v2` (default), `json` , `xml` , `text` , `x12` , `fhir` , `binary` , and `dicom` . `x12` rides any transport (see `samples/config/IB_PARTNER_X12.py`); `fhir` (ADR 0022) and `dicom` (ADR 0025) add their own on-demand codecs, and arbitrary bytes carry NUL-safely over the base64 `binary` path (ADR 0028). `FHIR()` is **outbound-only** (a FHIR REST *server* facade is not shipped). `DICOM()` runs in **both** directions from the one factory — an inbound C-STORE SCP listener and an outbound C-STORE SCU / C-ECHO sender to a downstream PACS ( `host` / `called_ae_title` configure the outbound peer).

- An **outbound FHIR() / Rest() destination to a SMART-secured server** (e.g. Epic, Oracle Health) can be wrapped with `with_smart_backend(...)` for OAuth2 client-credentials + signed-JWT authentication (ADR 0024). Import it as `from messagefoundry.transports.smart import with_smart_backend` — it is not re-exported from the top-level package.

The complete per-connector settings (TLS, retry, DoS guards, ACK mode, `simulate`, etc.) are documented in `CONNECTIONS.md`; each factory in `messagefoundry/config/wiring.py` is the **schema** for its transport.

## Bind an inbound to its Router

An inbound names its Router with the `router=` keyword ( `router="acme_adt_router"` above). The string must match a `@router` declared in some `.py` module loaded from the config dir; names resolve **globally** across the directory, so the inbound and its router can live in separate files. An inbound with no matching router fails `messagefoundry check`. (The router and handler are authored in code — covered in [Authoring Routers and Handlers](#).)

Validate the wiring before running it:

```
messagefoundry check --config samples/config --messages samples/messages/adt_a01.hl7
```

## Author a connection as data: `connections.toml`

A connection's transport config may instead live as **data** in an optional `connections.toml` next to the `.py` modules (ADR 0007). The loader merges these into the **same** registry the factories produce, so runtime, validation, and egress gating are identical. **Routing/transform logic stays code-first** — a data-authored inbound still binds a `router` declared in a `.py` module. Here is the shape, from `samples/config/connections.toml`:

```
[[inbound]]
name      = "IB_ACME_ADT_TCP"      # a second ACME intake, authored as data
transport = "mllp"
router    = "acme_adt_router"     # binds a router registered in a .py module
[inbound.settings]
port = 2700                        # inbound MLLP takes only a port
```

- The `transport` maps to the same factory ( `mllp` → `MLLP()` ), so a TOML connection produces a byte-identical spec and inherits every factory default and guard.
- **Secrets and per-environment peers use `{ env = "key" }`**, never an inline value (e.g. `host = { env = "acme_adt_host" }`, `port = { env = "acme_adt_port", cast = "int" }`). The file is repo-versionable and diffable.
- A name declared in **both** a `.py` module and `connections.toml` is a hard error (no silent shadowing).

Edit the file two ways, same file — by hand, or via the CLI (which is what the VS Code connection editor shells; it does a comment/format-preserving, validate-before-persist write):

```
messagefoundry connection list --config samples/config
messagefoundry connection upsert --config samples/config --data '{...}'
messagefoundry connection remove --config samples/config --name IB_ACME_ADT_TCP
```

`upsert` / `remove` validate the whole config dir (structure + connector build + the fail-closed `[egress]` allowlist) before persisting and roll back on failure.

## Try it end-to-end

Run the engine against the dev environment (the active `--env` is required), then send a synthetic message at the inbound's port:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
python samples/send_mllp.py samples/messages/adt_a01.hl7
```

Use only synthetic HL7 (as in `samples/messages/`) — never real PHI on a test feed.

## Authoring Routers and Handlers

A **Router** and a **Handler** are plain Python functions you write against the `messagefoundry` surface and register with the `@router` / `@handler` decorators. The Router sees *every* received message and decides which Handler(s) get it; each Handler filters, transforms the message, and returns `Send`s to outbound connections. Both are wired by name to a Connection — there is no enclosing "channel" object. The end-to-end template is `samples/results_relay/results_relay.py`; the simplest pair is `samples/config/adt.py`.

### 1. Write a Router (`@router`)

A Router takes the message and returns the **handler name(s)** to forward to — return `[]` to route nowhere (the message is still counted and logged `UNROUTED`, never dropped). It is the place to do fast, tolerant field peeks for routing decisions.

From `samples/config/adt.py`:

```
@router("adt_router")
def route(msg):
    if msg["MSH-9.1"] != "ADT":
        return [] # not ADT - routed nowhere (logged UNROUTED)
    return ["archive"]
```

The Router name ( `"adt_router"` ) is what an inbound Connection binds to: `inbound("IB_Test_ADT", MLLP(port=2575), router="adt_router")`. For the wiring conventions and per-connector settings, see `CONNECTIONS.md`.

## 2. Write a Handler ( `@handler` )

A Handler receives the message from a Router, then **filters** → **transforms** → **returns** `Send (s)`. Return `None` to filter the message out (logged `FILTERED`); return one `Send` or a list to fan out to multiple outbound connections.

From `samples/config/adt.py`:

```
@handler("archive")
def archive(msg):
    if msg["MSH-9.2"] not in EVENT_LABELS:
        return None # only admit/register/update events are archived (others FILTERED)
    mnemonic = FACILITY_MNEMONICS.get(msg["MSH-4"])
    if mnemonic:
        msg["MSH-4"] = mnemonic
    return Send("FILE-OUT_Test_ADT", msg)
```

The `Send` target ( `"FILE-OUT_Test_ADT"` ) names an `outbound(...)` Connection declared in the same config. To fan out, return a list — e.g. `samples/results_relay/results_relay.py` ends with `return [Send(OB_EHR, msg), Send(FILE_ARCHIVE, msg)]`.

## 3. The `Message` operations you'll use

Routers and Handlers work against the mutable HL7 `Message` in `messagefoundry/parsing/message.py` — never string-slice raw HL7; read/mutate through `Message` and re-encode. The methods you'll reach for (see the docstrings in that file for full signatures):

- **Peek a field** — `msg["PID-3"]` / `msg.field("OBX-3.1", occurrence=i)`; convenience properties `msg.message_code` (MSH-9.1), `msg.trigger_event` (MSH-9.2), `msg.control_id` (MSH-10).
- **Iterate repetitions / segments** — `msg.repetitions("PID-3")` walks a `~`-list; `msg.count_segments("OBX")` plus `field(occurrence=...)` walks repeating segments.
- **Mutate** — `msg["MSH-4"] = value` / `msg.set(path, value, occurrence=..., repetition=...)`; rebuild a repeating block with `msg.delete_segments("OBX")` + `msg.add_segment(line, index=...)` + `msg.add_repetition(...)`.
- **Read MSH separators** — never hardcode `|^~\&`. The repeating-segment rebuild in `samples/results_relay/results_relay.py` reads them from MSH-1/MSH-2 (its `_separators` helper) before joining components.
- **Re-encode** — `msg.encode()` (or just pass `msg` to a `Send`, which encodes for you).

A non-HL7 inbound ( `content_type` other than `hl7v2` ) delivers a `RawMessage` instead — read `.raw` / `.text` / `.json()` / `.xml()` (the XML accessor is XXE-safe via `defusedxml`: DOCTYPE, external-entity, and billion-laugh payloads raise) and `Send` a built string. For cross-field business-rule checks beyond what schema validation catches, compose the primitives in `parsing/consistency.py`, as `samples/consistency/validated_adt.py` does (raise `ConsistencyError` → dead-letter, or `return None` → filter). The three validation tiers are laid out in `HL7-VALIDATION.md`.

## 4. Translation tables (code sets)

A Router or Handler often maps a coded value to a downstream one — a sending-facility code to a mnemonic (the `FACILITY_MNEMONICS` lookup in the Handler above), an order code to a partner's code, a bed location to a room. Rather than hand-maintain a Python dict, you can back that lookup with a **translation table** (internally a *code set*): a `codesets/<name>.csv` file in your config dir (the name is the file stem). It **loads with the graph and reloads on promote**, and the lookup is **pure**, so it's safe under the staged pipeline.

```
codesets/facility_mnemonics.csv:
```

```
sending_facility,mnemonic
ACME_MAIN,ACMS
ACME_WEST,ACMW
```

Read it with `code_set("name")`, which returns a frozen, read-only mapping. Capture it once at a module's top level (or call it inline in a handler — both resolve):

```
from messagefoundry import code_set, handler, Send

FACILITIES = code_set("facility_mnemonics")    # loads with the graph; reloads on promote

@handler("archive")
def archive(msg):
    src = msg["MSH-4"]
    msg["MSH-4"] = FACILITIES.get(src, src)    # translate; pass through unchanged if unmapped
    return Send("FILE-OUT_Test_ADT", msg)
```

- **Missing key — you choose the behavior.** `code_set(...).get(key, default)` returns `default` on a miss (pass-through as above, or `""` to blank it); the subscript `code_set(...)[key]` **raises** on a miss, sending that message to its `ERROR` /dead-letter disposition (the strict "never deliver an unmapped value" path).
- **Single vs. multi-column.** One value column → the value is a scalar string; two or more → it's a `{header: cell}` dict ( `code_set("x")["k"]["mnemonic"]` ). Keys are exact-match, **case-sensitive** strings.

**Create and edit tables** by hand, or in the **Translation Tables** view of the VS Code extension — a grid editor (*New / Edit Translation Table*) that shells the offline `messagefoundry codeset` CLI (the validation authority):

```
messagefoundry codeset list    --config samples/config
messagefoundry codeset upsert  --config samples/config --data '{...}' # validate → atomic write
messagefoundry codeset rename  --config samples/config --name old --to new
messagefoundry codeset remove  --config samples/config --name old
```

A save is validated against the **same loader the engine uses** (no duplicate keys, no malformed file) and written atomically; a bad edit rolls back, and the change goes live through the usual **promote** (`POST /config/reload`).

**After a rename or remove, run `messagefoundry check`** — a handler's `code_set("old_name")` reference resolves at run time, so a plain `validate` won't catch a now-dangling name, but `check`'s dry-run will. Full reference: [CODESETS.md](#) and [CONFIGURATION.md](#); design record [ADR 0033](#). (For lookup data that lives in an **external** file or database rather than the bundle, see reference sets ([ADR 0006](#)) and the live `db_lookup` below.)

## 5. Purity rule (don't break this)

Routers and Handlers **must be pure**: message in → message(s) out, no external side effects. At-least-once delivery re-runs a transform after a crash and relies on the re-run producing identical output. Side effects (network, file, DB writes) belong in outbound Connections, not in your functions.

There are **two sanctioned exceptions**, both **read-only** and both available only inside a live Handler: a database read, `db_lookup(connection, statement, params)`, and a FHIR read/search, `fhir_lookup(connection, query)` — a read-by-id ( `"Patient/123"` ) or a search ( `"Patient?identifier=MRN|123"` ), GET-only. Either may return different data on a re-run, and that is accepted by design. Both are gated fail-closed — `db_lookup` by `[egress].allowed_db`, `fhir_lookup` by `[egress].allowed_http` — run off the event loop, and are **unavailable on a Router or in dry-run** (they raise). See [ADR 0010](#) and [ADR 0043](#).

## 6. The authoring dev loop

Run these from your config-repo root as you write — they touch no network and start no server.

```
# Static check: structural problems in the config dir
python -m messagefoundry validate --config samples/config

# See the wired Connection → Router → Handler → Connection graph
python -m messagefoundry graph --config samples/config

# Run real messages through Routers/Handlers WITHOUT sending — shows the disposition,
# selected handlers, and would-send payloads. Bodies are redacted unless you pass --show-phi.
# (--inbound selects which inbound to evaluate; it's required when the config wires more than one.)
python -m messagefoundry dryrun --config samples/config --messages samples/messages/adt_a01.hl7 --inbound
IB_Test_ADT

# The commit/CI gate: validate + dryrun (+ advisory ruff/mypy). Exit 0 only if every required check
# passes.
python -m messagefoundry check --config samples/config --messages samples/messages
```

`dryrun --show-phi` prints **full message bodies** (raw + would-send payloads) to stdout — that is PHI. Run it on synthetic HL7 only, and never redirect its output to a committed file, a ticket, or a CI log (PHI is redacted by default for exactly this reason). Wire `check` into your pre-commit hook / CI so a broken Router or Handler can't merge. For depth on connection settings see [CONNECTIONS.md](#); for the validation tiers see [HL7-VALIDATION.md](#).

---

## Operating with the console and the VS Code extension

MessageFoundry has two operator-facing UIs, and they do different jobs. The **web console** — the sole operator console, a browser UI the engine serves same-origin at `/ui` — monitors and operates a *running* engine over the localhost API: start/stop connections, browse the message log, watch health, manage users. The **VS Code extension** is for the *config author* — building and testing Connections/Routers/Handlers and promoting them to an engine. Neither touches the database directly; both go through the engine API. (The former PySide6 desktop console was retired — [BACKLOG #103](#); PySide6 now backs only the standalone test harness.)

## Opening and signing in to the console

The console is served by the engine itself, so start the engine first (note the active `--env` is required):

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

Then open the web console in a browser (the engine serves it at `/ui` by default, with the `messagefoundry-webconsole` distribution installed):

```
http://127.0.0.1:8765/ui
```

When the engine requires authentication (the default), a **Sign in** form appears first:

1. Enter your **username** and **password**, and pick a **Provider** — *Local* always; *Active Directory* appears only if the engine advertises AD.
2. If your account uses **two-factor (TOTP)**, you are prompted for the 6-digit code from your authenticator app (or a single-use recovery code) before you land on a page. An account with a **passkey** enrolled gets a *Use passkey* button here instead (ADR 0068; it needs the `[webauthn]` extra and a configured `public_origin`, and says so plainly when either is missing).
3. On a forced password change, the console chains a change-password step (local accounts only — AD passwords are changed in Active Directory).

On success the session rides an `HttpOnly mf_session` cookie, so navigating between pages doesn't re-prompt until the session expires or is revoked. The nav's **Account** menu leads to **My account**, which holds your **Password** (change it), **Multi-factor authentication** (enroll a TOTP authenticator — it shows a setup key + `otpauth://` URL for manual entry, then your one-time recovery codes — or turn it off), **Passkeys**, and **Active sessions** (inventory/revoke your own sessions); **Sign out** sits in the nav itself. Which pages and actions you can use is governed by RBAC — see [SECURITY.md](#) for roles and per-route permissions.

## A tour of the console pages

The pages hang off a **top nav** of hover/focus dropdowns: **Traffic** (Connections, Messages, Dead letters, Events), **Monitoring** (Status, Alerts, Flow & trends, Audit, Uploaded logs), **Admin** (Users, Configuration), and **Account** (My account, My security events). Every item is listed for everyone, but the *pages* are permission-gated server-side, so one you can't use refuses rather than hides. Flush right sit two live glyphs — an alerts bell and a health heart (green healthy, orange degraded, blinking red engine/DB stopped), repolled every ~15 s from every page — then **Sign out**.

- **Connections** — one row per endpoint (each inbound + each outbound). Tick rows and drive them from the dashboard toolbar's action dropdown: **Start / Stop / Restart** (either direction), **Reset stats**, and — for a *stopped and quiesced outbound* only — **Purge top** or **Purge all** (a step-up-unlocked, irreversible drain that cancels queued deliveries; it will not retry, and may be held for a second approver). Each row's name links to that connection's pre-filled **Messages** search, and an ⓘ link opens its read-only detail page. Columns are *Flag*



/ Connection / Dir / Status / In / Out / Queued / Errors / Alerts / Idle. A connection that failed to build or bind at startup shows a degraded **failed** status rather than taking the engine down (ADR 0031 — recovery is in [troubleshooting](#)).

- **Messages** — browse the message store (filter by connection, status, type, control ID, and a received-time window) and open one message at a time. Its detail page shows the metadata, the **Raw message** exactly as it arrived (preserved alongside the transformed form), **Deliveries** (per-destination status, attempts, last error), and **Events** (the audit trail), with a **Parse tree** → link to the structured HL7 view (server-parsed via the pure [parsing](#) library, so no body is ever re-rendered as markup). This is also where you **Replay** a message, or **Edit & resubmit** a copy of it. Viewing raw bodies is PHI access and is audited server-side with your username. A separate **Content search** page searches *inside* bodies by HL7 field path or substring — a bulk-PHI decrypt, so it is step-up gated and audited.
- **Status** — read-only health: engine (version, uptime, PID, inbounds running/total, endpoints in+out, engine-wide msg/s), store (path, size, free disk, journal mode, message/event/audit row counts), the effective security posture, and the active-passive **Cluster** roster + DR state. **Run integrity check** runs [PRAGMA quick\\_check](#) on demand, and **Reset statistics** zeroes the counters. When [\[service\].report\\_status](#) is on, a **Hosting service** badge reports the NSSM service's state — a read-out only, with no start/stop controls; manage the service itself per [SERVICE.md](#).
- **Users** (reading needs [users:read](#); every change needs [users:manage](#)) — RBAC administration across three cross-linked pages: **Users** (+ [New user](#)), then a per-user page with Profile, Roles, Channel scope, and account actions — *Reset password, Reset MFA, Sign out all sessions, Delete user*, **Roles** (user-definable custom roles over the built-in permission catalog, [ADR 0045](#)), and **AD group mappings**. Each body-carrying form opens inside a fresh step-up window, and every operation is audited server-side. Role definitions live in [SECURITY.md](#).
- **Alerts** — the engine's **active alert instances** (open and acknowledged) plus the loaded [\[alerts\]](#) rules ([ADR 0044](#) refining [ADR 0014](#)). Each instance carries severity, status, event type, connection, occurrence count, first/last seen, reason, and who acknowledged it, with **Ack / Resolve /** windowed **Suspend / Resume** actions. Rule *editing* stays config-file driven; the rule list is shown read-only (event type, connection, min depth, min age, severity, transports, cooldown — transports reported present-or-not, secrets omitted). The notifications themselves still fan out through the engine's AlertSink (see [Monitoring dispositions and troubleshooting](#)).

**Dead letters and replay.** The nav's **Dead letters** page lists dead-lettered deliveries newest-first — one row per message → destination that exhausted its retries (columns: *Failed / Channel / Destination / Type / Attempts / Last error / Message*) — with bulk **Replay all dead** — [<connection>](#) buttons, per-destination **Replay** [<connection>](#) → [<destination>](#) buttons, and one action to replay every dead delivery across all connections. Viewing the list needs [messages:read](#); replay needs [messages:replay](#) and is step-up (re-auth) gated server-side (and may be held for a second approver). Each row's **Message** link opens the audited detail page, where the **Deliveries** section shows the per-destination error and a per-message **Replay** — a second route to the same action. For diagnosing and clearing stuck deliveries, see the troubleshooting guidance in [EARLY-ADOPTER-GUIDE.md](#).



## The VS Code extension (for config authors)

The extension is a thin TypeScript UI that shells out to the `messagefoundry` CLI; it authors and tests interfaces, and can start/stop/restart a local engine and show its status — but **operating** and **monitoring** traffic is the web console's job. What it gives a config author:

- **Setup / scaffolding** — a **Home** launchpad with a **New Route Wizard** (Inbound → Router → Handler → Outbound generated as one module), **New Connection** (form → generates a `[TYPE]_[PARTNER]_[MESSAGE]` module like `IB_ACME_ADT`), **New Router/Handler**, **Generate Samples** (writes a synthetic, conformant corpus via `messagefoundry generate` — no PHI), and **Set Up Version Control & Checks** (puts the project under git with a `messagefoundry check` pre-commit hook).
- **Validate + graph** — *Validate on save* surfaces problems in the Problems panel; the **Components** view renders the wired graph (`messagefoundry graph`) by convention name, with **Filter** and **Group** controls and a row  **gear** to open a connection's `MLLP()` / `File()` settings in code.
- **Translation tables** — a **Translation Tables** view with a grid editor to create / edit / rename / delete a translation table (code set); it shells the `messagefoundry codeset` CLI (validate-on-save, atomic write) and offers **Promote** to apply. See [CODESETS.md](#).
- **Test Bench** — load `.hl7` files (each may hold many messages, split on `MSH`), **dry-run** them through the config without sending, and see each message's disposition, with a **Before/After** diff and a **Debug** step-through under `debugpy`. The load dialog opens to `messagefoundry.messageSetsDir` (default `samples/messages`).
- **Stage → Promote** — apply local config to a *running* engine: validate, pick a target environment, pre-flight a dry-run `POST /config/reload {dry_run:true}` against that target's `env()` values, confirm, then atomically swap the live graph. The engine requires auth, so the extension signs you in (token cached in VS Code SecretStorage).
- **AI assist** — an `@messagefoundry` chat participant ( `/explain` , `/transform` , `/router` , `/review` , `/migrate` , `/test` ) that is provider-agnostic and **only ever sends code + the config graph, never message bodies / PHI**.

Full feature and settings reference: [ide/README.md](#).

PHI note: `messagefoundry generate` and `dryrun` (which the Test Bench uses) print full message bodies to stdout/stderr — run them only against synthetic HL7 (e.g. `samples/messages/adt_a01.hl7`), and never redirect their output to a committed file, ticket, or CI log.

---

## Monitoring dispositions and troubleshooting

MessageFoundry never accepts-and-drops a message: **every message a connection receives is persisted and counted before it is ACKed**, and its *disposition* is recorded as it flows through the pipeline. This section is how you watch those dispositions, recover failed deliveries, and get paged when a lane stalls. For the underlying guarantees see [ARCHITECTURE.md](#) and the staged-pipeline rationale in [ADR 0001](#).

## What each disposition means

A message's status moves through the **staged pipeline** ( `ingress` -> `routed` -> `outbound` ). The store's finalizer is the **single authority** that sets the final disposition — it only finalizes once every handler's work resolves, so one delivered handler can't mark a message done while a sibling is still in flight:

| Status                 | What it means to you                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| RECEIVED               | Persisted at ingress and ACKed. The count is booked; routing hasn't run yet.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| ROUTED                 | The Router selected at least one Handler; the message is awaiting transform/delivery.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| UNROUTED               | The Router ran but chose <b>no</b> Handler. Not an error — logged, kept, not delivered. Check the Router logic if you expected a destination.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| FILTERED               | Every Handler ran but delivered nothing (filtered out). Expected for drop-by-design rules; surprising drops mean a Handler filter is too aggressive.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| NOT_DEPLOYED           | Every destination the Handlers addressed is present in the config but marked <code>deployed = false</code> , so the engine declined the send and queued nothing (ADR 0111). Not an error, and deliberately <i>not</i> <code>FILTERED</code> — the Handler did choose to send; the engine declined the destination. Each declined leg also lands as a <code>not_deployed</code> event on the message, and that event is in the compliance floor, so it survives even <code>[diagnostics].message_events = "off"</code> . A message that still reached a <code>deployed</code> sibling therefore finalizes <code>PROCESSED</code> , with the skipped leg on the record. To bring the destination up, set <code>deployed = true</code> , supply that connection's <code>env()</code> values, and reload — start/restart from the console is refused ( <code>409</code> ), because deploying is a config change, not a runtime action. See <a href="#">CONNECTIONS.md</a> . |
| PROCESSED              | Every selected Handler transformed and <b>all</b> destinations delivered. The happy path.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| ERROR /<br>dead-letter | A stage failed. A decode/parse/strict-validate failure is recorded <i>before</i> ingress (and NAK'd — see below); a routing/transform/delivery failure is recorded <i>after</i> the ACK as <code>ERROR</code> plus an <code>AlertSink</code> event. A delivery that exhausts retries becomes a <b>dead-letter</b> you can inspect and replay.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |

The key operator shift under the staged pipeline: **an AA ACK means "received and persisted," not "delivered."** A post-ingress failure is a disposition + alert, not a NAK.

## Where to watch dispositions

- **Web console -> Messages.** The message browser has a **status** filter — type a disposition (e.g. `unrouted`, `error`) to narrow the list, then open a message to see its raw body, parse tree, deliveries, and audit trail. On the Connections page, a connection's name links straight to Messages pre-filtered to that connection.
- **Web console -> Connections.** The dashboard shows each connection's live status plus per-connection counts, including an **errored** column, so a climbing error count on one feed is visible at a glance.
- **API.** `GET /messages?status=error` (and `&channel_id=`, `&message_type=`) is the filter the console uses; `GET /stats` returns outbox-by-status + in-pipeline depth; `GET /status` returns engine uptime, running/stopped channel counts, and DB size/free-disk; `GET /metrics` exposes a Prometheus exposition (aggregate counts/latency keyed by connection + status, no PHI) for an external scraper — it works on a base install, and the `[otel]` extra adds only the optional OpenTelemetry/OTLP export seam. All require the `monitoring:read` (stats/status/metrics) or `messages:read` (the Messages browser) permission — see [SECURITY.md](#).

## The ERROR / dead-letter path: inspect and replay

A delivery dead-letters when its retries are exhausted. Retry behavior is per-outbound (defaults in `[delivery]` — see `CONFIGURATION.md`):

- `retry_max_attempts` **unset = retry forever** (the conservative default; under FIFO the failing head blocks its lane until it succeeds or is purged). Set a finite value to opt into retry-then-dead-letter.
- A partner **AR reject fails fast** (no retry); an **AE NAK / transient transport failure is retried** with backoff.

To recover:

1. **Find the dead-letters.** Console: open the **Dead letters** page (or the message itself from **Messages**) and read its delivery row's **Last error**. API: `GET /dead-letters` (optionally `?channel_id=&destination_name=`) lists dead deliveries newest-first; each row carries `last_error`.
2. **Fix the cause** (the downstream endpoint, the transform, the config).
3. **Replay.** `POST /dead-letters/replay` re-queues the dead deliveries (optionally scoped by `channel_id` / `destination_name`); each affected message reverts from `error` to `received` and re-drains. Already-delivered rows are left alone. Replay requires the `messages:replay` permission and is **step-up (re-auth) gated**, and may be held for a second approver when `[approvals]` is configured. (In the console, the **Dead letters** page lists these and offers the per-connection, per-destination, and replay-everything buttons directly; this API path is the equivalent for scripting and automation.)

Replaying re-transmits real message bodies — it is audited per acting user. Treat it like any PHI action (`PHI.md`).

## Alerting: fire-and-forward notifications

The engine raises operational alert events — `connection_stopped` (a lane halted by the `stop` internal-error policy), `queue_buildup` (a backlog past its depth/age threshold), `storage_threshold` (the store grew past `[retention].max_db_mb`), and `cert_expiry` (a monitored TLS certificate nearing expiry) — through an `AlertSink`. With no `[alerts]` transport configured these are just logged at `WARNING`; configure a transport and they fan out to it (`alert_sinks.py`).

Configure alerts in the `[alerts]` section (`CONFIGURATION.md`):

```
[alerts]
webhook_url = "https://hooks.example.com/mf"    # POSTs each event as JSON (Slack/Teams/PagerDuty)
email_smtp_host = "smtp.example.com"
email_from = "messagefoundry@example.com"
email_to = ["oncall@example.com"]
# Optional per-event routing/severity/suppression rules (first match wins) — ADR 0014:
[[alerts.rules]]
event_type = "connection_stopped"
severity = "critical"
transports = ["webhook"]
```

The SMTP password is a secret — supply it via `MEFOR_ALERTS_EMAIL_PASSWORD`, never the file. Per-event severity, transport routing, thresholds, suppression, and cooldown are tuned with ordered `[[alerts.rules]]` tables (ADR 0014); an event matching no rule notifies every configured transport at `warning`.

**Important:** the *notifications* are **fire-and-forward** — once webhook/email delivery is attempted there is no send log, so for the notification itself you rely on your webhook/email target. The engine does, however, keep **alert state**: an alert instance per `(event type, connection)` that you can query with `GET /alerts/active` and act on with `POST /alerts/{id}/ack · /resolve · /suspend · /resume`, which is what the console's **Alerts** page drives (ADR 0044). `GET /alerts/rules` remains a read-only view of the *loaded* transport config and rule set (secrets and recipients omitted). Instances and payloads carry only the connection name and queue shape, never message content (no PHI).

## Common problems

- **Sender got a NAK (AE/AR).** A decode/parse/strict-validate failure rejects *synchronously* at the listener and records **ERROR before** any ingress row — the message never entered the pipeline. Fix the inbound HL7 (or relax `validation.strict` on that connection). Treat the message body as untrusted data, not a malformed instruction.
- **Sender got AA but nothing was delivered.** Expected under ACK-on-receipt: routing/transform/delivery failures happen *after* the ACK. Look at the message's disposition ( `UNROUTED` / `FILTERED` / `NOT_DEPLOYED` / `ERROR` ) and the AlertSink — **not** the ACK — for the outcome.
- **A lane stopped processing.** A `connection_stopped` alert means an outbound's worker halted on an internal/code error ( `internal_error = stop` ). The messages are preserved for replay; fix the cause, then reload/restart the connection.
- **A connection shows failed.** A connection that can't build or bind **at startup** (bad settings, a port already in use) is isolated as a degraded `failed` status instead of taking the engine down — every other lane keeps running (ADR 0031). Fix the config/bind, then recover it: restart an inbound ( `POST /connections/{name}/start` ), or reload to rebuild a failed outbound. (Reload itself stays fail-fast — a broken config is rejected whole, never partially applied.)
- **Backlog growing.** A `queue_buildup` alert usually means a retry-forever head is blocking its FIFO lane, or the downstream is down. Check the destination, then inspect/purge or replay the blocking row.
- **Console can't reach the engine.** The API binds `127.0.0.1:8765` by default and requires auth; confirm the engine is serving ( `python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev` ), that the `messagefoundry-webconsole` distribution is installed and the console has not been turned off ( `[security].serve_web_console` ), and that your browser is pointed at that host/port's `/ui`.
- **Low disk / store growing.** `GET /status` reports DB size and free disk; a `storage_threshold` alert fires past `[retention].max_db_mb`. Tune retention in `[retention]` (CONFIGURATION.md) — purges null PHI bodies while keeping the message/disposition rows, so counts and audit stay intact. The row is kept; its PHI columns — operator-attached `metadata` included — are blanked.

---

## Where to go next

- 
- **Concepts in depth** — [MENTAL-MODEL.md](#), [ARCHITECTURE.md](#) (diagrams: [architecture-diagram.md](#))
  - **Narrative onboarding, install to production** — [EARLY-ADOPTER-GUIDE.md](#), [INSTALL-GUIDE.md](#), [SYSTEM-REQUIREMENTS.md](#)
  - **Connections reference** — [CONNECTIONS.md](#), [ADR 0007](#) ( [connections.toml](#) )
  - **Translation tables (code sets)** — [CODESETS.md](#), [ADR 0033](#)
  - **Service settings & environments** — [CONFIGURATION.md](#)
  - **Validation tiers** — [HL7-VALIDATION.md](#)
  - **Run as a service** — [SERVICE.md](#)
  - **Security, RBAC, TLS** — [SECURITY.md](#), [DEPLOYMENT.md](#)
  - **PHI handling & encryption-at-rest** — [PHI.md](#)
  - **VS Code extension** — [ide/README.md](#)
  - **What's built vs. planned** — [FEATURE-MAP.md](#), [README.md](#)
  - **Design records** — [ADR 0001](#) (staged pipeline), [ADR 0004](#) (payload-agnostic ingress), [ADR 0010](#) ( [db\\_lookup](#) ), [ADR 0012](#) (X12), [ADR 0014](#) (alerting rules)

---

© 2026 **MEFOR-ORG**, a non-profit · MessageFoundry is licensed under **AGPL-3.0-or-later** · Document generated July 2026 for MessageFoundry v0.3.2.

MessageFoundry and the MessageFoundry word mark are trademarks of MEFOR-ORG. This document describes an **Early Access** release; check [pypi.org/project/messagefoundry](https://pypi.org/project/messagefoundry) for the current version. Verify specifics against your own deployment and the engine's reference documentation.