
MessageFoundry User Guide

Early Access · MessageFoundry v0.3.2 · July 2026

MessageFoundry is an open-source, Python interface engine for healthcare — a modern alternative to Mirth Connect and Corepoint. It **receives, routes, transforms, and validates** messages between systems (HL7 v2.x by default, payload-agnostic for other formats), with routing and handling expressed as **ordinary Python you own and version-control**. This guide is **task-oriented** ("how do I..."): it walks you from a clean machine to a running engine, then through authoring Connections/Routers/Handlers and operating and monitoring the system. It links to the reference docs rather than restating them.

Contents

- What MessageFoundry is, and how to use this guide
- Installing and running the engine
- Quickstart: send your first message
- Authoring Connections
- Authoring Routers and Handlers
- Operating with the console and the VS Code extension
- Monitoring dispositions and troubleshooting
- Where to go next

What MessageFoundry is, and how to use this guide

MessageFoundry receives, routes, transforms, and validates messages between systems: **HL7 v2.x by default**, and payload-agnostic for other formats (JSON, XML/SOAP, X12 EDI, database records). Unlike a legacy engine's embedded scripting language or a locked-in low/no-code GUI, the routing and handling logic is **ordinary Python you own and version-control** — and connection setup in particular can be pure data (a TOML file, edited by hand or in a VS Code GUI). The engine runs headless as an asyncio service; you operate it from a browser **web console** (served same-origin at `/ui`) and a VS Code extension, both talking to it over a localhost HTTP/WebSocket API.

The mental model in one read

The configuration is a **graph wired by name** — there is **no "channel" object** that bundles a source, filters, transforms, and destinations the way a legacy engine does. Instead you wire four building blocks, using the project's exact vocabulary:

- **Connection** — an endpoint that **receives** (`inbound()`) or **sends** (`outbound()`) messages: MLLP, TCP, File, REST, SOAP, Database, SFTP/FTP. Every message in or out is counted and logged.
- **Router** (`@router`) — a pure Python function bound to **one** inbound Connection. It sees every received message and returns the **name(s)** of the Handler(s) to forward to (or none, to filter).
- **Handler** (`@handler`) — a pure Python function that takes a message, **filters** → **transforms**, then returns `Send(...)` s naming one or more outbound Connections.
- **Message store** — durable persistence + the staged queue. SQLite (WAL) by default; PostgreSQL or SQL Server for production.

The edges between these are just names resolved at config load: an inbound names its Router, a Router returns Handler names, a Handler `Send` s to a named outbound. To understand a feed, follow the names; to change it, change a name. (See [samples/config/IB_ACME_ADT.py](#) and [samples/config/adt.py](#) for a complete, runnable example of this wiring.)

Under the hood, a received message flows through a **staged pipeline** of three persisted stages — **ingress** (the raw message, committed *before* the ACK) → **routed** (one row per Handler the Router selected) → **outbound** (one row per destination) — each drained by its own asyncio worker. Because nothing is ever silently dropped, every message carries a **disposition** that the store records as it flows: `RECEIVED` → `ROUTED` / `UNROUTED` → `PROCESSED` / `FILTERED` / `NOT_DEPLOYED` / `ERROR` . This is the count-and-log promise; operators read the disposition (and alerts), not the ACK, to confirm a message reached its destination. The full reliability model (at-least-once delivery, why routers/handlers must be pure) is in [ADR 0001](#).

Who this guide is for

This guide serves two audiences: **operators** who install, run, and monitor the engine, and **config authors** (integration developers and analysts) who wire Connections, Routers, and Handlers. It links out to the reference docs rather than restating them. If you want a lighter, narrative onboarding first, read [EARLY-ADOPTER-GUIDE.md](#) (install-to-production rollout) or [MENTAL-MODEL.md](#) (the concepts above, in depth).

Where the depth lives — the reference map:

When you want...	Read
The full architecture, modularity standard, and dependency rules	ARCHITECTURE.md (diagrams: architecture-diagram.md)
Connection types, settings, and the graph (incl. <code>connections.toml</code>)	CONNECTIONS.md , ADR 0007
Translation tables (code sets) — the grid editor + <code>codeset</code> CLI	CODESETS.md , ADR 0033
Service settings, environments, and <code>env()</code> values	CONFIGURATION.md
Running as a Windows service (NSSM)	SERVICE.md
Auth, RBAC, audit, and TLS	SECURITY.md , DEPLOYMENT.md
PHI handling and encryption-at-rest	PHI.md

When you want...	Read
The staged pipeline / reliability; payload-agnostic ingress; the read-only <code>db_lookup</code> ; X12	ADR 0001, ADR 0004, ADR 0010, ADR 0012
What's built vs. planned	FEATURE-MAP.md, README.md

A note on PHI before you run anything: this engine carries PHI, and a few commands (notably `dryrun` and `generate`) print **full message bodies** to stdout. Use **synthetic HL7 only** in examples, and never redirect that output to a committed file, ticket, or CI log. To build *realistic* test data from real traffic without ever handling PHI, use the built de-identification framework via the standalone tee relay's `python -m tee anonymize-captures` (fail-closed; ADR 0030) rather than hand-writing synthetic messages. See [PHI.md](#) for the hard rules.

How the rest of this guide is laid out

The remaining sections are a **path**, in the order you'll actually work through them:

1. **Install** the engine and scaffold your config repo.
2. Send your **first message** end to end (e.g. `samples/messages/adt_a01.hl7` via `samples/send_mllp.py`) against the bundled sample config.
3. **Author Connections** (in Python or `connections.toml`).
4. **Author Routers and Handlers** to route and transform.
5. **Operate** via the browser web console (`/ui` , `messagefoundry-webconsole`) and the VS Code extension (`ide/`).
6. **Monitor and troubleshoot** with dispositions, alerts, dead-letter triage, and replay.

Work through them in order the first time; afterward, jump to the task you need.

Installing and running the engine

This section takes you from a clean machine to a running engine and an attached console. It covers the **developer / checkout** path (running the bundled `samples/config`) and points at the **pinned-wheel** consumer path where they differ.

Two install styles exist. To **try MessageFoundry against the sample config** (this guide's running example), install from a checkout (`pip install -e .`). To **deploy your own interfaces**, install the pinned wheel and scaffold a config repo with `messagefoundry init` — the full consumer model is [INSTALL-GUIDE.md](#).

1. Check prerequisites

- **Python 3.14+** (64-bit). Everything else (the SQLite store, NSSM for the service) is auto-provisioned or bundled.
- **git** if you are cloning the checkout or standing up a config repo.
- **Administrator/elevation** only for the Windows service step (6).

Full hardware, OS, database, and port details: [SYSTEM-REQUIREMENTS.md](#).

2. Install the engine

Create a virtual environment and install. For the **checkout / contributor** path (gives you `samples/config`, `samples/messages/`, and `samples/send_mllp.py`):

```
python -m venv .venv
.\.venv\Scripts\Activate.ps1      # Linux/macOS: . .venv/bin/activate
pip install -e .                  # core runtime + SQLite store
```

For a **deployment**, pin the published wheel instead (and verify its provenance — see [INSTALL-GUIDE.md](#)):

```
pip install "messagefoundry==0.3.2"
```

Add only the extras a host actually needs (each is opt-in and lazy-imported):

```
pip install -e packaging/messagefoundry-webconsole # the browser web console (/ui) – the operator UI (step 5)
pip install -e ".[harness]"                        # the standalone PySide6 test harness GUI
pip install -e ".[postgres]"                       # PostgreSQL store backend
pip install -e ".[sqlserver]"                      # SQL Server store backend – also needs the OS-level Microsoft ODBC Driver 18
pip install -e ".[sftp]"                          # SFTP transport for the REMOTEFILER connector
pip install -e ".[fhir]"                          # FHIR codec (R5/R4B/STU3) + FHIR() outbound (ADR 0022)
pip install -e ".[dicom]"                         # DICOM C-STORE SCP + codec – headers/SR only, no pixels (ADR 0025)
pip install -e ".[webauthn]"                      # browser WebAuthn passkeys for the /ui console (ADR 0068)
pip install -e ".[otel]"                          # OpenTelemetry/OTLP export seam (the /metrics endpoint itself needs no extra)
```

(For a deployment wheel, the same extras apply: `pip install "messagefoundry[harness]==0.3.2"`, and the web console installs as its own wheel `pip install "messagefoundry-webconsole==0.2.15"`.) SQLite is the zero-dependency default — you need no extra to run the sample config.

3. Run the engine headless (dev)

From the repo root, against the bundled sample config:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

- `--config` points at the directory of Connection/Router/Handler modules (here `samples/config/`, which includes `IB_ACME_ADT.py` and its transform `adt.py`).
- `--db` is the SQLite message store path (created on first run).
- `--env` is required — `serve` refuses to start without it. The active environment is a free-form **name** (`dev` / `staging` / `prod`, or a custom name) that does two things: it selects the value file `environments/<env>.toml` that `env("...")` lookups resolve against, and it sets the instance's **PHI posture**

(`data_class` / `production`). Built-in names carry a default posture; a custom name must declare it. See [CONFIGURATION.md](#).

When the engine runs from somewhere other than the repo root (e.g. under the service), anchor the value files with `--project-root <repo-root>` so `env()` values don't silently resolve empty — see [INSTALL-GUIDE.md](#).

Network / auth posture. The API binds `127.0.0.1:8765` and **requires authentication** by default. A non-loopback bind without TLS is refused at startup; configure native TLS (or an upstream terminator) to expose it. Details: [SECURITY.md](#) and [DEPLOYMENT.md](#).

Store encryption (PHI instances). On a PHI instance (`[security].handles_real_patient_data = true`), `serve` **refuses to start** — in every environment, `dev` and `staging` included, not just production — if no store encryption key is configured. Mint one with `messagefoundry gen-key` (set it as `MEFOR_STORE_ENCRYPTION_KEY`), or on Windows DPAPI-protect it to a file with `messagefoundry protect-key --generate --out <file>` and point `[store].encryption_key_file` at it. The full key story is in [PHI.md](#).

Confirm it's up:

```
curl http://127.0.0.1:8765/health
```

4. Scaffold your own config repo

When you're ready to author real interfaces, scaffold a standalone, `check` -green config repo instead of editing the samples:

```
messagefoundry init ./my-config-repo
cd my-config-repo
```

It writes a runnable starter feed under `config/`, `environments/dev.toml` + `prod.toml`, a synthetic fixture under `messages/sets/`, a `messagefoundry.toml`, a `requirements.txt` pinning the engine, and a CI `check.yml`. Validate and run it:

```
pip install -r requirements.txt
messagefoundry check --config config --messages messages/sets
messagefoundry serve --config config --env dev
```

Put it under version control with **Set Up Version Control & Checks** in the IDE (or a plain `git init`). During setup you choose **where the repo is stored** — *on this machine only* (fine for a single-box, non-HA engine or local dev) or on a *shared remote* (recommended for HA, a team, or off-machine backup); change it any time with **Config Repo Storage Location**. Choosing local vs. remote storage, and why secrets/PHI never land in the repo, are covered in [VERSION-CONTROL.md](#) (with the full deployment model in [INSTALL-GUIDE.md](#)).

5. Open the admin console (in a browser)

The console is the **browser web console** served same-origin by the engine at `/ui` — install the `messagefoundry-webconsole` wheel alongside the engine and it is served automatically, **on by default** at a loopback bind (turn it off with `[security].serve_web_console = false`). With the engine running, browse to:

```
http://127.0.0.1:8765/ui
```

The web console prompts for sign-in (authentication is on by default). Source: `packaging/messagefoundry-webconsole/`. (The former PySide6 desktop console was retired — BACKLOG #103; PySide6 now backs only the standalone test harness.)

6. Run as a Windows service (NSSM)

For production on Windows, run the engine as a background service via **NSSM** — it starts on boot, restarts on crash, captures stdout/stderr to rotating logs, and stops with Ctrl+C so connections drain cleanly. From an **elevated** PowerShell:

```
.\scripts\service\install-service.ps1 -Environment prod
```

`-Environment` is **required** (it becomes `serve --env`, just like step 3). The script is idempotent (re-run to reconfigure), auto-downloads a SHA-256-pinned NSSM if one isn't on `PATH`, and defaults to service name `MessageFoundry`, config `<repo>\samples\config`, store + logs under `C:\ProgramData\MessageFoundry`, bind `127.0.0.1:8765`. Override paths/port/account with flags, e.g.:

```
.\scripts\service\install-service.ps1 -Environment prod -Port 9000 `
  -Config D:\hl7\config -DataDir D:\MessageFoundry `
  -ServiceAccount "NT SERVICE\MessageFoundry" # least-privilege; auto-grants the needed ACLs
```

Manage and remove it:

```
nssm start MessageFoundry
nssm status MessageFoundry
nssm stop MessageFoundry
.\scripts\service\uninstall-service.ps1 # elevated; leaves logs + store in place
```

The complete procedure — least-privilege accounts, locking down the config/log directories, DPAPI key protection, update-vs-reinstall, and troubleshooting — is in `SERVICE.md`.

A note on PHI-emitting commands

`messagefoundry dryrun` and `messagefoundry generate` print **full message bodies** to stdout/stderr (`dryrun` only with `--show-phi`; redacted otherwise). Run them against **synthetic HL7 only** — never real PHI — and never redirect their output into a committed file, ticket, or CI log. See `PHI.md`.

Quickstart: send your first message

This walkthrough uses only the shipped samples in `samples/config/` — no editing required. You'll start the engine, push a synthetic HL7 ADT message over MLLP, and watch it get received, routed, and archived to a file.

The sample inbound that does the work is `IB_Test_ADT` in `samples/config/adt.py`: an MLLP listener on **port 2575** whose Router forwards `ADT` messages to the `archive` Handler, which writes them to `./out/adt/{MSH-10}.hl7` via the `FILE-OUT_Test_ADT` outbound connection. (The config dir also wires other sample feeds — ACME ADT on 2600, X12, immunizations — but this Quickstart only exercises `IB_Test_ADT`.)

1. Start the engine

In one terminal, from the repo root, run the engine against the sample config. The active `--env` is **required** — use `dev` for local work:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

This loads the config modules, opens (or creates) the SQLite store at `./messagefoundry.db`, serves the localhost API, and starts every sample listener. The startup log announces the active environment and posture. Leave it running. The API binds `127.0.0.1` by default; first-run auth/console details are covered in the install guide — see `INSTALL-GUIDE.md`.

2. Send the sample message

In a **second** terminal (also from the repo root), send the shipped synthetic ADT^A01 over MLLP with the helper in `samples/send_mllp.py`:

```
python samples/send_mllp.py samples/messages/adt_a01.hl7
```

The script defaults to `--host 127.0.0.1 --port 2575`, which matches `IB_Test_ADT`, so no flags are needed. To be explicit (or to target a different listener), pass them:

```
python samples/send_mllp.py samples/messages/adt_a01.hl7 --host 127.0.0.1 --port 2575
```

The file it sends, `samples/messages/adt_a01.hl7`, is synthetic — never substitute real PHI here.

3. What to expect

`send_mllp.py` reuses the engine's MLLP framing, waits for the framed ACK, and prints it:

```
--- ACK ---
MSH|^~\&|...|...|...|ACK^A01|...|P|2.5.1
MSA|AA|MSG00001
...
```

An `MSA|AA` is a positive acknowledgement (`AA` = Application Accept). Under the staged pipeline the ACK means **received and durably persisted** (status `RECEIVED` at the ingress stage), *not* final delivery — see the ACK-on-receipt model in [ARCHITECTURE.md](#). A moment later the message is routed and delivered: because `adt_a01.h17` is an `ADT^A01`, the Router forwards it to the `archive` Handler, the Handler stamps a downstream facility mnemonic *if one is mapped for the sending facility* (none is for this sample, so its MSH-4 is left unchanged) and sends it on, and the file outbound writes it under `./out/adt/`. Its disposition advances `RECEIVED` → `ROUTED` → `PROCESSED`. (A non-`ADT` message would be logged `UNROUTED`; an `ADT` event the Handler drops would be `FILTERED` — nothing is silently discarded.)

Confirm the delivered file landed (its name is the message's MSH-10 control ID, `MSG00001`):

```
ls ./out/adt/
```

4. Where to see it land

Two complementary ways to inspect what happened:

- **The web console Messages page.** Open the web console at `/ui` in a browser, then use **Traffic** → **Messages** to find the message, its disposition, the original raw body, and the per-destination delivery. The web console talks only to the engine's API — see [packaging/messagefoundry-webconsole/](#).
- **A dryrun preview (no engine needed).** To see exactly how the config *would* route and transform a message without sending it anywhere, run:

```
python -m messagefoundry dryrun --config samples/config --messages samples/messages/adt_a01.h17 --inbound IB_Test_ADT
```

Dryrun prints the resolved inbound, disposition, selected handlers, and would-send payloads (`--inbound` names which inbound to evaluate — required here because the sample config wires several). **Its output is PHI-bearing** — message bodies are redacted by default and only included with `--show-phi`. The samples here are synthetic, but never pipe `dryrun` (or `generate`) output to a committed file or a CI log.

Now change what it does

Once the round trip works, make it yours: edit the Router and Handler in `samples/config/adt.py` to change routing and transforms (see [Authoring Routers and Handlers](#)), add or retarget connections in `samples/config/connections.toml` — by hand or via the VS Code editor ([ADR 0007](#)) — and follow the connection naming convention in [CONNECTIONS.md](#). Run `python -m messagefoundry check --config samples/config --messages samples/messages` before committing config changes.

Authoring Connections

A **Connection** is an endpoint that either *receives* messages (an **inbound** source) or *sends* them (an **outbound** destination). MLLP and File are the two most common transports, both shipped today (plus raw TCP, X12, REST, SOAP, Database, and SFTP/FTP — see the full catalog and per-setting reference in [CONNECTIONS.md](#)).

Connections carry only *transport* config; routing/transform *logic* lives in code-first Routers and Handlers (see [Authoring Routers and Handlers](#)).

You author a connection one of two ways — **as code** (a `.py` module) or **as data** (`connections.toml`). Both desugar into the same registry, so they coexist freely.

Name your connection

Use the convention `[TYPE]_[PARTNER]_[MESSAGE]`:

- **TYPE** — transport + direction code: `IB` / `OB` (inbound/outbound MLLP), `FILE-IN` / `FILE-OUT`, `TCP-IN` / `TCP-OUT`, `X12-IN` / `X12-OUT`, etc. (the full table is in `CONNECTIONS.md`).
- **PARTNER** — the system on the other end (`ACME`, `Epic`, `Test`).
- **MESSAGE** — the HL7 message code (`ADT`, `ORU`, `VXU`, ...) or `MIXED` / `ALL`.

Example: `IB_ACME_ADT` = inbound MLLP from ACME carrying ADT. Names are plain strings, so hyphens and mixed case (`FILE-OUT_Test_ADT`) are fine. Router/Handler names are *not* connections and don't follow this formula.

Author a code-first inbound and outbound

In a config module, call the `inbound()` / `outbound()` factories with a transport spec (`MLLP()`, `File()`, ...). Here is the shape, drawn from `samples/config/IB_ACME_ADT.py`:

```
from messagefoundry import MLLP, env, inbound, outbound

inbound("IB_ACME_ADT", MLLP(port=2600), router="acme_adt_router")
outbound("OB_ACME_ADT", MLLP(host=env("acme_adt_host"), port=env("acme_adt_port", cast=int)))
```

Key points the sample demonstrates:

- **Inbound MLLP takes only a port** — passing `host` is a wiring error. The listen interface is the service-level `[inbound].bind_host` (loopback in DEV, a specific NIC in PROD), an operator setting, not authored here.
- **Outbound MLLP needs a host and port**. Anything that differs by environment (a downstream peer, a credential) uses `env("key")`, resolved per instance from `environments/<env>.toml` (and `MEFOR_VALUE_<KEY>` for secrets) — so one module runs unchanged in every environment. A referenced-but-undefined value fails loud at load, never a silent blank host.
- For a **File** endpoint, use `File(directory="./out/adt")` (in) / `File(directory=..., filename="{MSH-10}.h17")` (out). For non-HL7 bodies, set the inbound's `content_type` so the body routes as a `RawMessage` instead of being HL7-parsed — the shipped set is `hl7v2` (default), `json`, `xml`, `text`, `x12`, `fhir`, `binary`, and `dicom`. `x12` rides any transport (see `samples/config/IB_PARTNER_X12.py`); `fhir` (ADR 0022) and `dicom` (ADR 0025) add their own on-demand codecs, and arbitrary bytes carry NUL-safely over the base64 `binary` path (ADR 0028). Two of the newer connectors are direction-specific: `FHIR()` is **outbound-only** and the `DICOM()` C-STORE listener is **inbound-only**.
- An **outbound FHIR() / Rest()** destination to a **SMART-secured server** (e.g. Epic, Oracle Health) can be wrapped with `with_smart_backend(...)` for OAuth2 client-credentials + signed-JWT authentication (ADR 0024). Import it as `from messagefoundry.transports.smart import with_smart_backend` — it is not re-

exported from the top-level package.

The complete per-connector settings (TLS, retry, DoS guards, ACK mode, `simulate`, etc.) are documented in `CONNECTIONS.md`; each factory in `messagefoundry/config/wiring.py` is the **schema** for its transport.

Bind an inbound to its Router

An inbound names its Router with the `router=` keyword (`router="acme_adt_router"` above). The string must match a `@router` declared in some `.py` module loaded from the config dir; names resolve **globally** across the directory, so the inbound and its router can live in separate files. An inbound with no matching router fails `messagefoundry check`. (The router and handler are authored in code — covered in [Authoring Routers and Handlers](#).)

Validate the wiring before running it:

```
messagefoundry check --config samples/config --messages samples/messages/adt_a01.h17
```

Author a connection as data: `connections.toml`

A connection's transport config may instead live as **data** in an optional `connections.toml` next to the `.py` modules ([ADR 0007](#)). The loader merges these into the **same** registry the factories produce, so runtime, validation, and egress gating are identical. **Routing/transform logic stays code-first** — a data-authored inbound still binds a `router` declared in a `.py` module. Here is the shape, from `samples/config/connections.toml`:

```
[[inbound]]
name      = "IB_ACME_ADT_TCP"      # a second ACME intake, authored as data
transport = "mllp"
router    = "acme_adt_router"     # binds a router registered in a .py module
[inbound.settings]
port      = 2700                  # inbound MLLP takes only a port
```

- The `transport` maps to the same factory (`mllp` → `MLLP()`), so a TOML connection produces a byte-identical spec and inherits every factory default and guard.
- **Secrets and per-environment peers use `{ env = "key" }`**, never an inline value (e.g. `host = { env = "acme_adt_host" }`, `port = { env = "acme_adt_port", cast = "int" }`). The file is repo-versionable and diffable.
- A name declared in **both** a `.py` module and `connections.toml` is a hard error (no silent shadowing).

Edit the file two ways, same file — by hand, or via the CLI (which is what the VS Code connection editor shells; it does a comment/format-preserving, validate-before-persist write):

```
messagefoundry connection list --config samples/config
messagefoundry connection upsert --config samples/config --data '{...}'
messagefoundry connection remove --config samples/config --name IB_ACME_ADT_TCP
```

`upsert` / `remove` validate the whole config dir (structure + connector build + the fail-closed `[egress]` allowlist) before persisting and roll back on failure.

Try it end-to-end

Run the engine against the dev environment (the active `--env` is required), then send a synthetic message at the inbound's port:

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
python samples/send_mllp.py samples/messages/adt_a01.hl7
```

Use only synthetic HL7 (as in `samples/messages/`) — never real PHI on a test feed.

Authoring Routers and Handlers

A **Router** and a **Handler** are plain Python functions you write against the `messagefoundry` surface and register with the `@router` / `@handler` decorators. The Router sees *every* received message and decides which Handler(s) get it; each Handler filters, transforms the message, and returns `Send`s to outbound connections. Both are wired by name to a Connection — there is no enclosing "channel" object. The end-to-end template is `samples/results_relay/results_relay.py`; the simplest pair is `samples/config/adt.py`.

1. Write a Router (`@router`)

A Router takes the message and returns the **handler name(s)** to forward to — return `[]` to route nowhere (the message is still counted and logged `UNROUTED`, never dropped). It is the place to do fast, tolerant field peeks for routing decisions.

From `samples/config/adt.py`:

```
@router("adt_router")
def route(msg):
    if msg["MSH-9.1"] != "ADT":
        return [] # not ADT - routed nowhere (logged UNROUTED)
    return ["archive"]
```

The Router name (`"adt_router"`) is what an inbound Connection binds to: `inbound("IB_Test_ADT", MLLP(port=2575), router="adt_router")`. For the wiring conventions and per-connector settings, see `CONNECTIONS.md`.

2. Write a Handler (`@handler`)

A Handler receives the message from a Router, then **filters** → **transforms** → **returns** `Send (s)`. Return `None` to filter the message out (logged `FILTERED`); return one `Send` or a list to fan out to multiple outbound connections.

From `samples/config/adt.py`:

```

@handler("archive")
def archive(msg):
    if msg["MSH-9.2"] not in EVENT_LABELS:
        return None # only admit/register/update events are archived (others FILTERED)
    mnemonic = FACILITY_MNEMONICS.get(msg["MSH-4"])
    if mnemonic:
        msg["MSH-4"] = mnemonic
    return Send("FILE-OUT_Test_ADT", msg)

```

The `Send` target (`"FILE-OUT_Test_ADT"`) names an `outbound(...)` Connection declared in the same config. To fan out, return a list — e.g. `samples/results_relay/results_relay.py` ends with `return [Send(OB_EHR, msg), Send(FILE_ARCHIVE, msg)]`.

3. The **Message** operations you'll use

Routers and Handlers work against the mutable HL7 `Message` in `messagefoundry/parsing/message.py` — never string-slice raw HL7; read/mutate through `Message` and re-encode. The methods you'll reach for (see the docstrings in that file for full signatures):

- **Peek a field** — `msg["PID-3"]` / `msg.field("OBX-3.1", occurrence=i)`; convenience properties `msg.message_code` (MSH-9.1), `msg.trigger_event` (MSH-9.2), `msg.control_id` (MSH-10).
- **Iterate repetitions / segments** — `msg.repetitions("PID-3")` walks a `~`-list; `msg.count_segments("OBX")` plus `field(occurrence=...)` walks repeating segments.
- **Mutate** — `msg["MSH-4"] = value` / `msg.set(path, value, occurrence=..., repetition=...)`; rebuild a repeating block with `msg.delete_segments("OBX")` + `msg.add_segment(line, index=...)` + `msg.add_repetition(...)`.
- **Read MSH separators** — never hardcode `|^~\&`. The repeating-segment rebuild in `samples/results_relay/results_relay.py` reads them from MSH-1/MSH-2 (its `_separators` helper) before joining components.
- **Re-encode** — `msg.encode()` (or just pass `msg` to a `Send`, which encodes for you).

A non-HL7 inbound (`content_type` other than `hl7v2`) delivers a `RawMessage` instead — read `.raw` / `.text` / `.json()` / `.xml()` (the XML accessor is XXE-safe via `defusedxml`: DOCTYPE, external-entity, and billion-laughs payloads raise) and `Send` a built string. For cross-field business-rule checks beyond what schema validation catches, compose the primitives in `parsing/consistency.py`, as `samples/consistency/validated_adt.py` does (raise `ConsistencyError` → dead-letter, or `return None` → filter). The three validation tiers are laid out in `HL7-VALIDATION.md`.

4. Translation tables (code sets)

A Router or Handler often maps a coded value to a downstream one — a sending-facility code to a mnemonic (the `FACILITY_MNEMONICS` lookup in the Handler above), an order code to a partner's code, a bed location to a room. Rather than hand-maintain a Python dict, you can back that lookup with a **translation table** (internally a *code set*): a `codesets/<name>.csv` file in your config dir (the name is the file stem). It **loads with the graph and reloads on promote**, and the lookup is **pure**, so it's safe under the staged pipeline.

codesets/facility_mnemonics.csv:

```
sending_facility,mnemonic
ACME_MAIN,ACMS
ACME_WEST,ACMW
```

Read it with `code_set("name")`, which returns a frozen, read-only mapping. Capture it once at a module's top level (or call it inline in a handler — both resolve):

```
from messagefoundry import code_set, handler, Send

FACILITIES = code_set("facility_mnemonics")    # loads with the graph; reloads on promote

@handler("archive")
def archive(msg):
    src = msg["MSH-4"]
    msg["MSH-4"] = FACILITIES.get(src, src)    # translate; pass through unchanged if unmapped
    return Send("FILE-OUT_Test_ADT", msg)
```

- **Missing key — you choose the behavior.** `code_set(...).get(key, default)` returns `default` on a miss (pass-through as above, or `""` to blank it); the subscript `code_set(...)[key]` **raises** on a miss, sending that message to its `ERROR` /dead-letter disposition (the strict "never deliver an unmapped value" path).
- **Single vs. multi-column.** One value column → the value is a scalar string; two or more → it's a `{header: cell}` dict (`code_set("x")["k"]["mnemonic"]`). Keys are exact-match, **case-sensitive** strings.

Create and edit tables by hand, or in the **Translation Tables** view of the VS Code extension — a grid editor (*New / Edit Translation Table*) that shells the offline `messagefoundry codeset` CLI (the validation authority):

```
messagefoundry codeset list    --config samples/config
messagefoundry codeset upsert  --config samples/config --data '{...}' # validate → atomic write
messagefoundry codeset rename  --config samples/config --name old --to new
messagefoundry codeset remove  --config samples/config --name old
```

A save is validated against the **same loader the engine uses** (no duplicate keys, no malformed file) and written atomically; a bad edit rolls back, and the change goes live through the usual **promote** (`POST /config/reload`).

After a rename or remove, run `messagefoundry check` — a handler's `code_set("old_name")` reference resolves at run time, so a plain `validate` won't catch a now-dangling name, but `check`'s dry-run will. Full reference: [CODESETS.md](#) and [CONFIGURATION.md](#); design record [ADR 0033](#). (For lookup data that lives in an **external** file or database rather than the bundle, see reference sets ([ADR 0006](#)) and the live `db_lookup` below.)

5. Purity rule (don't break this)

Routers and Handlers **must be pure**: message in → message(s) out, no external side effects. At-least-once delivery re-runs a transform after a crash and relies on the re-run producing identical output. Side effects (network, file, DB writes) belong in outbound Connections, not in your functions.

The **one sanctioned exception** is a Handler making a **live, read-only** `db_lookup(connection, statement, params)` for enrichment/gating — its result may differ on a re-run, and that is accepted by design. It is read-only, gated by `[egress].allowed_db`, runs off the event loop, and is **unavailable on a Router or in dry-run** (it raises). See [ADR 0010](#).

6. The authoring dev loop

Run these from your config-repo root as you write — they touch no network and start no server.

```
# Static check: structural problems in the config dir
python -m messagefoundry validate --config samples/config

# See the wired Connection → Router → Handler → Connection graph
python -m messagefoundry graph --config samples/config

# Run real messages through Routers/Handlers WITHOUT sending — shows the disposition,
# selected handlers, and would-send payloads. Bodies are redacted unless you pass --show-phi.
# (--inbound selects which inbound to evaluate; it's required when the config wires more than one.)
python -m messagefoundry dryrun --config samples/config --messages samples/messages/adt_a01.hl7 --inbound
IB_Test_ADT

# The commit/CI gate: validate + dryrun (+ advisory ruff/mypy). Exit 0 only if every required check
# passes.
python -m messagefoundry check --config samples/config --messages samples/messages
```

`dryrun --show-phi` prints **full message bodies** (raw + would-send payloads) to stdout — that is PHI. Run it on synthetic HL7 only, and never redirect its output to a committed file, a ticket, or a CI log (PHI is redacted by default for exactly this reason). Wire `check` into your pre-commit hook / CI so a broken Router or Handler can't merge. For depth on connection settings see [CONNECTIONS.md](#); for the validation tiers see [HL7-VALIDATION.md](#).

Operating with the console and the VS Code extension

MessageFoundry has two operator-facing UIs, and they do different jobs. The **web console** — the sole operator console, a browser UI the engine serves same-origin at `/ui` — monitors and operates a *running* engine over the localhost API: start/stop connections, browse the message log, watch health, manage users. The **VS Code extension** is for the *config author* — building and testing Connections/Routers/Handlers and promoting them to an engine. Neither touches the database directly; both go through the engine API. (The former PySide6 desktop console was retired — [BACKLOG #103](#); PySide6 now backs only the standalone test harness.)

Opening and signing in to the console

The console is served by the engine itself, so start the engine first (note the active `--env` is required):

```
python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev
```

Then open the web console in a browser (the engine serves it at `/ui` by default, with the `messagefoundry-webconsole` distribution installed):

http://127.0.0.1:8765/ui

When the engine requires authentication (the default), a **Sign in** form appears first:

1. Enter your **username** and **password**, and pick a **Provider** — *Local* always; *Active Directory* appears only if the engine advertises AD.
2. If your account uses **two-factor (TOTP)**, you are prompted for the 6-digit code from your authenticator app (or a single-use recovery code) before you land on a page. An account with a **passkey** enrolled gets a *Use passkey* button here instead (ADR 0068; it needs the `[webauthn]` extra and a configured `public_origin`, and says so plainly when either is missing).
3. On a forced password change, the console chains a change-password step (local accounts only — AD passwords are changed in Active Directory).

On success the session rides an `HttpOnly mf_session` cookie, so navigating between pages doesn't re-prompt until the session expires or is revoked. The nav's **Account** menu leads to **My account**, which holds your **Password** (change it), **Multi-factor authentication** (enroll a TOTP authenticator — it shows a setup key + `otpauth://` URL for manual entry, then your one-time recovery codes — or turn it off), **Passkeys**, and **Active sessions** (inventory/revoke your own sessions); **Sign out** sits in the nav itself. Which pages and actions you can use is governed by RBAC — see [SECURITY.md](#) for roles and per-route permissions.

A tour of the console pages

The pages hang off a **top nav** of hover/focus dropdowns: **Traffic** (Connections, Messages, Dead letters, Events), **Monitoring** (Status, Alerts, Flow & trends, Audit, Uploaded logs), **Admin** (Users, Configuration), and **Account** (My account, My security events). Every item is listed for everyone, but the *pages* are permission-gated server-side, so one you can't use refuses rather than hides. Flush right sit two live glyphs — an alerts bell and a health heart (green healthy, orange degraded, blinking red engine/DB stopped), repolled every ~15 s from every page — then **Sign out**.

- **Connections** — one row per endpoint (each inbound + each outbound). Tick rows and drive them from the dashboard toolbar's action dropdown: **Start / Stop / Restart** (either direction), **Reset stats**, and — for a *stopped and quiesced outbound* only — **Purge top** or **Purge all** (a step-up-unlocked, irreversible drain that cancels queued deliveries; it will not retry, and may be held for a second approver). Each row's name links to that connection's pre-filled **Messages** search, and an ⓘ link opens its read-only detail page. Columns are *Flag / Connection / Dir / Status / In / Out / Queued / Errors / Alerts / Idle*. A connection that failed to build or bind at startup shows a degraded **failed** status rather than taking the engine down (ADR 0031 — recovery is in [troubleshooting](#)).
- **Messages** — browse the message store (filter by connection, status, type, control ID, and a received-time window) and open one message at a time. Its detail page shows the metadata, the **Raw message** exactly as it arrived (preserved alongside the transformed form), **Deliveries** (per-destination status, attempts, last error), and **Events** (the audit trail), with a **Parse tree** → link to the structured HL7 view (server-parsed via the pure `parsing` library, so no body is ever re-rendered as markup). This is also where you **Replay** a message, or **Edit**

& resubmit a copy of it. Viewing raw bodies is PHI access and is audited server-side with your username. A separate **Content search** page searches *inside* bodies by HL7 field path or substring — a bulk-PHI decrypt, so it is step-up gated and audited.

- **Status** — read-only health: engine (version, uptime, PID, inbounds running/total, endpoints in+out, engine-wide msg/s), store (path, size, free disk, journal mode, message/event/audit row counts), the effective security posture, and the active-passive **Cluster** roster + DR state. **Run integrity check** runs `PRAGMA quick_check` on demand, and **Reset statistics** zeroes the counters. When `[service].report_status` is on, a **Hosting service** badge reports the NSSM service's state — a read-out only, with no start/stop controls; manage the service itself per `SERVICE.md`.
- **Users** (reading needs `users:read`; every change needs `users:manage`) — RBAC administration across three cross-linked pages: **Users** (+ `New user`, then a per-user page with Profile, Roles, Channel scope, and account actions — *Reset password*, *Reset MFA*, *Sign out all sessions*, *Delete user*), **Roles** (user-definable custom roles over the built-in permission catalog, `ADR 0045`), and **AD group mappings**. Each body-carrying form opens inside a fresh step-up window, and every operation is audited server-side. Role definitions live in `SECURITY.md`.
- **Alerts** — the engine's **active alert instances** (open and acknowledged) plus the loaded `[alerts]` rules (`ADR 0044` refining `ADR 0014`). Each instance carries severity, status, event type, connection, occurrence count, first/last seen, reason, and who acknowledged it, with **Ack** / **Resolve** / windowed **Suspend** / **Resume** actions. Rule *editing* stays config-file driven; the rule list is shown read-only (event type, connection, min depth, min age, severity, transports, cooldown — transports reported present-or-not, secrets omitted). The notifications themselves still fan out through the engine's AlertSink (see `Monitoring dispositions and troubleshooting`).

Dead letters and replay. The nav's **Dead letters** page lists dead-lettered deliveries newest-first — one row per message → destination that exhausted its retries (columns: *Failed* / *Channel* / *Destination* / *Type* / *Attempts* / *Last error* / *Message*) — with bulk **Replay all dead** — `<connection>` buttons, per-destination **Replay** `<connection>` → `<destination>` buttons, and one action to replay every dead delivery across all connections. Viewing the list needs `messages:read`; replay needs `messages:replay` and is step-up (re-auth) gated server-side (and may be held for a second approver). Each row's **Message** link opens the audited detail page, where the **Deliveries** section shows the per-destination error and a per-message **Replay** — a second route to the same action. For diagnosing and clearing stuck deliveries, see the troubleshooting guidance in `EARLY-ADOPTER-GUIDE.md`.

The VS Code extension (for config authors)

The extension is a thin TypeScript UI that shells out to the `messagefoundry` CLI; it authors and tests interfaces, and can start/stop/restart a local engine and show its status — but **operating** and **monitoring** traffic is the web console's job. What it gives a config author:

- **Setup / scaffolding** — a **Home** launchpad with a **New Route Wizard** (Inbound → Router → Handler → Outbound generated as one module), **New Connection** (form → generates a `[TYPE]_[PARTNER]_[MESSAGE]` module like `IB_ACME_ADT`), **New Router/Handler**, **Generate Samples** (writes a synthetic, conformant corpus via `messagefoundry generate` — no PHI), and **Set Up Version Control & Checks** (puts the project under git with a `messagefoundry check` pre-commit hook).

- **Validate + graph** — *Validate on save* surfaces problems in the Problems panel; the **Components** view renders the wired graph (`messagefoundry graph`) by convention name, with **Filter** and **Group** controls and a row  **gear** to open a connection's `MLLP()` / `File()` settings in code.
- **Translation tables** — a **Translation Tables** view with a grid editor to create / edit / rename / delete a translation table (code set); it shells the `messagefoundry codeset` CLI (validate-on-save, atomic write) and offers **Promote** to apply. See [CODESETS.md](#).
- **Test Bench** — load `.hl7` files (each may hold many messages, split on `MSH`), **dry-run** them through the config without sending, and see each message's disposition, with a **Before/After** diff and a **Debug** step-through under `debugpy`. The load dialog opens to `messagefoundry.messageSetsDir` (default `samples/messages`).
- **Stage → Promote** — apply local config to a *running* engine: validate, pick a target environment, pre-flight a dry-run `POST /config/reload {dry_run:true}` against that target's `env()` values, confirm, then atomically swap the live graph. The engine requires auth, so the extension signs you in (token cached in VS Code SecretStorage).
- **AI assist** — an `@messagefoundry` chat participant (`/explain`, `/transform`, `/router`, `/review`, `/migrate`, `/test`) that is provider-agnostic and **only ever sends code + the config graph, never message bodies / PHI**.

Full feature and settings reference: [ide/README.md](#).

PHI note: `messagefoundry generate` and `dryrun` (which the Test Bench uses) print full message bodies to stdout/stderr — run them only against synthetic HL7 (e.g. `samples/messages/adt_a01.hl7`), and never redirect their output to a committed file, ticket, or CI log.

Monitoring dispositions and troubleshooting

MessageFoundry never accepts-and-drops a message: **every message a connection receives is persisted and counted before it is ACKed**, and its *disposition* is recorded as it flows through the pipeline. This section is how you watch those dispositions, recover failed deliveries, and get paged when a lane stalls. For the underlying guarantees see [ARCHITECTURE.md](#) and the staged-pipeline rationale in [ADR 0001](#).

What each disposition means

A message's status moves through the **staged pipeline** (`ingress` -> `routed` -> `outbound`). The store's finalizer is the **single authority** that sets the final disposition — it only finalizes once every handler's work resolves, so one delivered handler can't mark a message done while a sibling is still in flight:

Status	What it means to you
RECEIVED	Persisted at ingress and ACKed. The count is booked; routing hasn't run yet.
ROUTED	The Router selected at least one Handler; the message is awaiting transform/delivery.

Status	What it means to you
UNROUTED	The Router ran but chose no Handler. Not an error — logged, kept, not delivered. Check the Router logic if you expected a destination.
FILTERED	Every Handler ran but delivered nothing (filtered out). Expected for drop-by-design rules; surprising drops mean a Handler filter is too aggressive.
NOT_DEPLOYED	Every destination the Handlers addressed is present in the config but marked <code>deployed = false</code> , so the engine declined the send and queued nothing (ADR 0111). Not an error, and deliberately <i>not</i> <code>FILTERED</code> — the Handler did choose to send; the engine declined the destination. Each declined leg also lands as a <code>not_deployed</code> event on the message, and that event is in the compliance floor, so it survives even <code>[diagnostics].message_events = "off"</code> . A message that still reached a <i>deployed</i> sibling therefore finalizes <code>PROCESSED</code> , with the skipped leg on the record. To bring the destination up, set <code>deployed = true</code> , supply that connection's <code>env()</code> values, and reload — start/restart from the console is refused (409), because deploying is a config change, not a runtime action. See CONNECTIONS.md .
PROCESSED	Every selected Handler transformed and all destinations delivered. The happy path.
ERROR / dead-letter	A stage failed. A decode/parse/strict-validate failure is recorded <i>before</i> ingress (and NAK'd — see below); a routing/transform/delivery failure is recorded <i>after</i> the ACK as <code>ERROR</code> plus an <code>AlertSink</code> event. A delivery that exhausts retries becomes a dead-letter you can inspect and replay.

The key operator shift under the staged pipeline: **an AA ACK means "received and persisted," not "delivered."** A post-ingress failure is a disposition + alert, not a NAK.

Where to watch dispositions

- **Web console -> Messages.** The message browser has a **status** filter — type a disposition (e.g. `unrouted`, `error`) to narrow the list, then open a message to see its raw body, parse tree, deliveries, and audit trail. On the Connections page, a connection's name links straight to Messages pre-filtered to that connection.
- **Web console -> Connections.** The dashboard shows each connection's live status plus per-connection counts, including an **errored** column, so a climbing error count on one feed is visible at a glance.
- **API.** `GET /messages?status=error` (and `&channel_id=`, `&message_type=`) is the filter the console uses; `GET /stats` returns outbox-by-status + in-pipeline depth; `GET /status` returns engine uptime, running/stopped channel counts, and DB size/free-disk; `GET /metrics` exposes a Prometheus exposition (aggregate counts/latency keyed by connection + status, no PHI) for an external scraper — it works on a base install, and the `[otel]` extra adds only the optional OpenTelemetry/OTLP export seam. All require the `monitoring:read` (stats/status/metrics) or `messages:read` (the Messages browser) permission — see [SECURITY.md](#).

The ERROR / dead-letter path: inspect and replay

A delivery dead-letters when its retries are exhausted. Retry behavior is per-outbound (defaults in `[delivery]` — see [CONFIGURATION.md](#)):

- `retry_max_attempts` **unset = retry forever** (the conservative default; under FIFO the failing head blocks its lane until it succeeds or is purged). Set a finite value to opt into retry-then-dead-letter.
- A partner **AR reject fails fast** (no retry); an **AE NAK / transient transport failure is retried** with backoff.

To recover:

1. **Find the dead-letters.** Console: open the **Dead letters** page (or the message itself from **Messages**) and read its delivery row's **Last error**. API: `GET /dead-letters` (optionally `?channel_id=&destination_name=`) lists dead deliveries newest-first; each row carries `last_error`.
2. **Fix the cause** (the downstream endpoint, the transform, the config).
3. **Replay.** `POST /dead-letters/replay` re-queues the dead deliveries (optionally scoped by `channel_id` / `destination_name`); each affected message reverts from `error` to `received` and re-drains. Already-delivered rows are left alone. Replay requires the `messages:replay` permission and is **step-up (re-auth) gated**, and may be held for a second approver when `[approvals]` is configured. (In the console, the **Dead letters** page lists these and offers the per-connection, per-destination, and replay-everything buttons directly; this API path is the equivalent for scripting and automation.)

Replaying re-transmits real message bodies — it is audited per acting user. Treat it like any PHI action (PHI.md).

Alerting: fire-and-forward notifications

The engine raises operational alert events — `connection_stopped` (a lane halted by the `stop` internal-error policy), `queue_buildup` (a backlog past its depth/age threshold), `storage_threshold` (the store grew past `[retention].max_db_mb`), and `cert_expiry` (a monitored TLS certificate nearing expiry) — through an `AlertSink`. With no `[alerts]` transport configured these are just logged at `WARNING`; configure a transport and they fan out to it (`alert_sinks.py`).

Configure alerts in the `[alerts]` section (`CONFIGURATION.md`):

```
[alerts]
webhook_url = "https://hooks.example.com/mf" # POSTs each event as JSON (Slack/Teams/PagerDuty)
email_smtp_host = "smtp.example.com"
email_from = "messagefoundry@example.com"
email_to = ["oncall@example.com"]
# Optional per-event routing/severity/suppression rules (first match wins) — ADR 0014:
[[alerts.rules]]
event_type = "connection_stopped"
severity = "critical"
transports = ["webhook"]
```

The SMTP password is a secret — supply it via `MEFOR_ALERTS_EMAIL_PASSWORD`, never the file. Per-event severity, transport routing, thresholds, suppression, and cooldown are tuned with ordered `[[alerts.rules]]` tables (ADR 0014); an event matching no rule notifies every configured transport at `warning`.

Important: the *notifications* are **fire-and-forward** — once webhook/email delivery is attempted there is no send log, so for the notification itself you rely on your webhook/email target. The engine does, however, keep **alert state**: an alert instance per `(event_type, connection)` that you can query with `GET /alerts/active` and act on with `POST /alerts/{id}/ack` · `/resolve` · `/suspend` · `/resume`, which is what the console's **Alerts** page drives

(ADR 0044). `GET /alerts/rules` remains a read-only view of the *loaded* transport config and rule set (secrets and recipients omitted). Instances and payloads carry only the connection name and queue shape, never message content (no PHI).

Common problems

- **Sender got a NAK (AE/AR).** A decode/parse/strict-validate failure rejects *synchronously* at the listener and records `ERROR` **before** any ingress row — the message never entered the pipeline. Fix the inbound HL7 (or relax `validation.strict` on that connection). Treat the message body as untrusted data, not a malformed instruction.
- **Sender got AA but nothing was delivered.** Expected under ACK-on-receipt: routing/transform/delivery failures happen *after* the ACK. Look at the message's disposition (`UNROUTED` / `FILTERED` / `NOT_DEPLOYED` / `ERROR`) and the AlertSink — **not** the ACK — for the outcome.
- **A lane stopped processing.** A `connection_stopped` alert means an outbound's worker halted on an internal/code error (`internal_error = stop`). The messages are preserved for replay; fix the cause, then reload/restart the connection.
- **A connection shows `failed`.** A connection that can't build or bind **at startup** (bad settings, a port already in use) is isolated as a degraded `failed` status instead of taking the engine down — every other lane keeps running (ADR 0031). Fix the config/bind, then recover it: restart an inbound (`POST /connections/{name}/start`), or reload to rebuild a failed outbound. (Reload itself stays fail-fast — a broken config is rejected whole, never partially applied.)
- **Backlog growing.** A `queue_buildup` alert usually means a retry-forever head is blocking its FIFO lane, or the downstream is down. Check the destination, then inspect/purge or replay the blocking row.
- **Console can't reach the engine.** The API binds `127.0.0.1:8765` by default and requires auth; confirm the engine is serving (`python -m messagefoundry serve --config samples/config --db ./messagefoundry.db --env dev`), that the `messagefoundry-webconsole` distribution is installed and the console has not been turned off (`[security].serve_web_console`), and that your browser is pointed at that host/port's `/ui`.
- **Low disk / store growing.** `GET /status` reports DB size and free disk; a `storage_threshold` alert fires past `[retention].max_db_mb`. Tune retention in `[retention]` (CONFIGURATION.md) — purges null PHI bodies while keeping the message/disposition rows, so counts and audit stay intact. The row is kept; its PHI columns — operator-attached `metadata` included — are blanked.

Where to go next

- **Concepts in depth** — `MENTAL-MODEL.md`, `ARCHITECTURE.md` (diagrams: `architecture-diagram.md`)
- **Narrative onboarding, install to production** — `EARLY-ADOPTER-GUIDE.md`, `INSTALL-GUIDE.md`, `SYSTEM-REQUIREMENTS.md`
- **Connections reference** — `CONNECTIONS.md`, ADR 0007 (`connections.toml`)
- **Translation tables (code sets)** — `CODESETS.md`, ADR 0033
- **Service settings & environments** — `CONFIGURATION.md`

- **Validation tiers** — [HL7-VALIDATION.md](#)
- **Run as a service** — [SERVICE.md](#)
- **Security, RBAC, TLS** — [SECURITY.md](#), [DEPLOYMENT.md](#)
- **PHI handling & encryption-at-rest** — [PHI.md](#)
- **VS Code extension** — [ide/README.md](#)
- **What's built vs. planned** — [FEATURE-MAP.md](#), [README.md](#)
- **Design records** — [ADR 0001](#) (staged pipeline), [ADR 0004](#) (payload-agnostic ingress), [ADR 0010](#) ([db_lookup](#)), [ADR 0012](#) (X12), [ADR 0014](#) (alerting rules)

© 2026 **MEFOR-ORG**, a non-profit · MessageFoundry is licensed under **AGPL-3.0-or-later** · Document generated July 2026 for MessageFoundry v0.3.2.

MessageFoundry and the MessageFoundry word mark are trademarks of MEFOR-ORG. This document describes an **Early Access** release; check pypi.org/project/messagefoundry for the current version. Verify specifics against your own deployment and the engine's reference documentation.